

# NYU FRE 7773 - Week 2

---

*Machine Learning in Financial Engineering*

Ethan Rosenthal

# Linear Models for Regression

---

*Machine Learning in Financial Engineering*

Ethan Rosenthal

Computer Setup?

# Linear Models for Regression

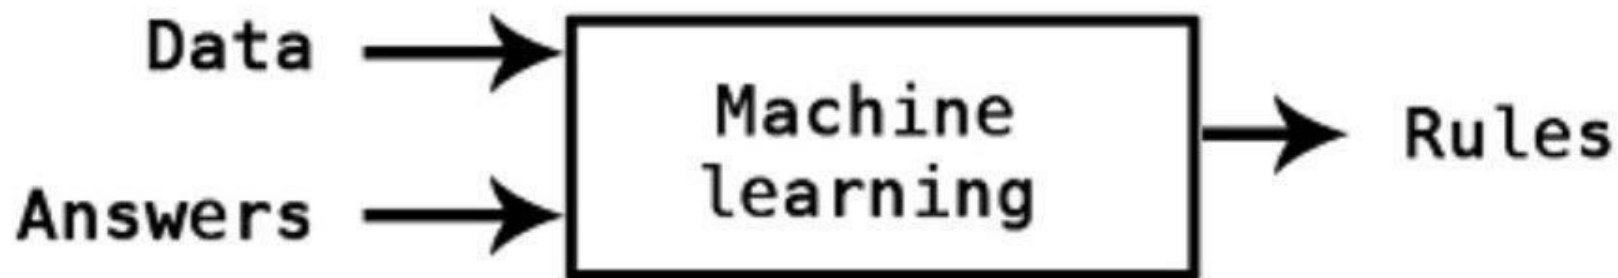
---

*Machine Learning in Financial Engineering*

Ethan Rosenthal

# Linear Models for Regression

- **Regression:** for when we are predicting continuous values.
- **Linear Model:** a model that is linear in the model parameters/coefficients.



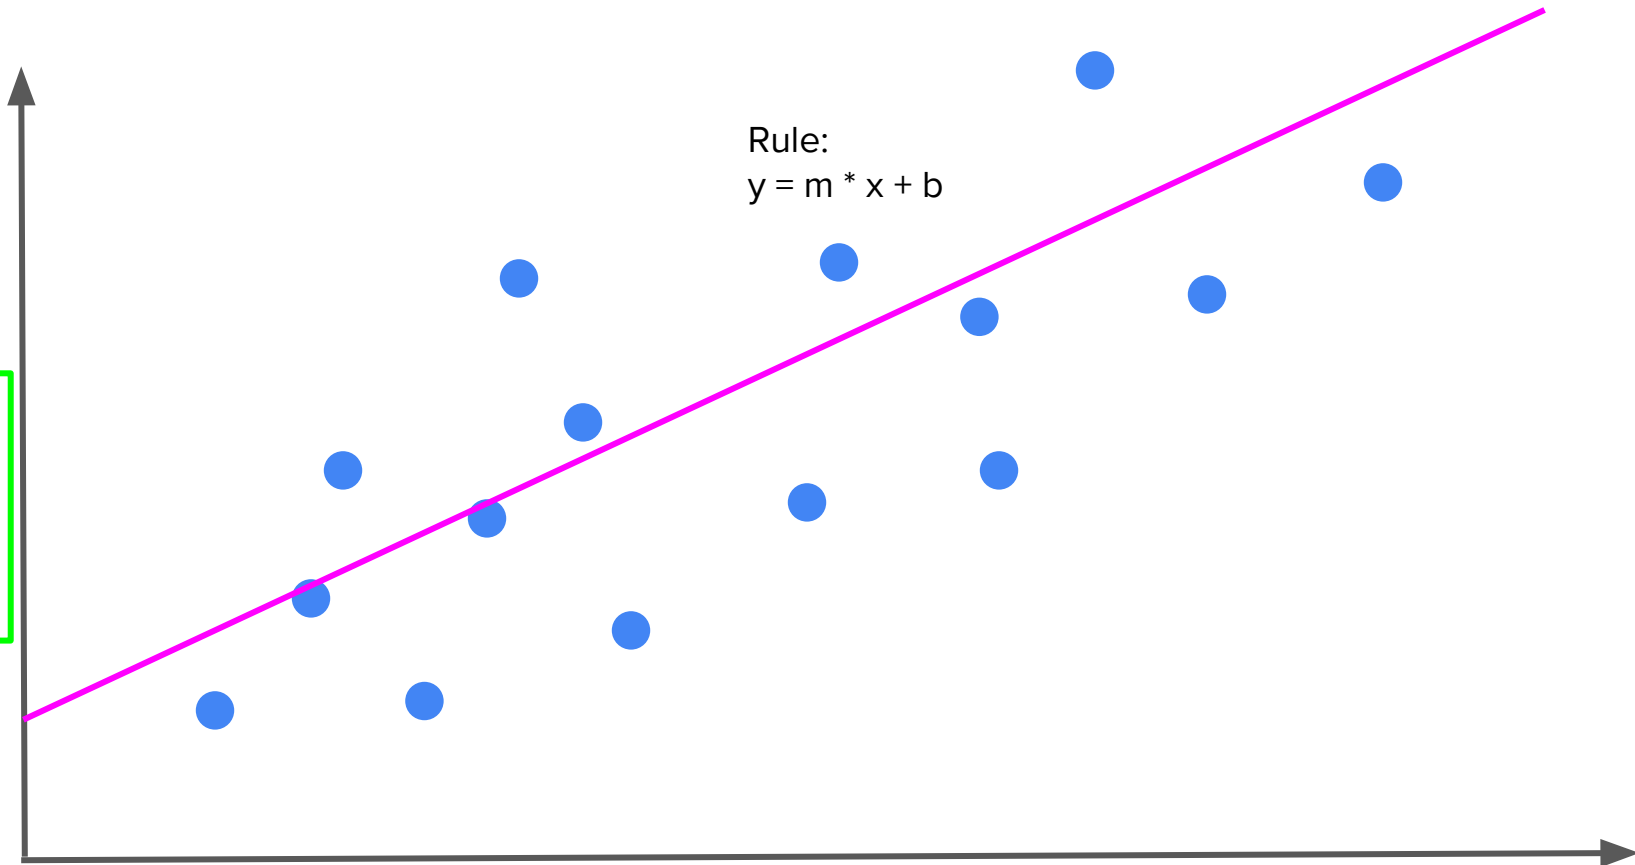
Rule:  
 $y = m * x + b$

List Price

Square Footage

Data

Answers



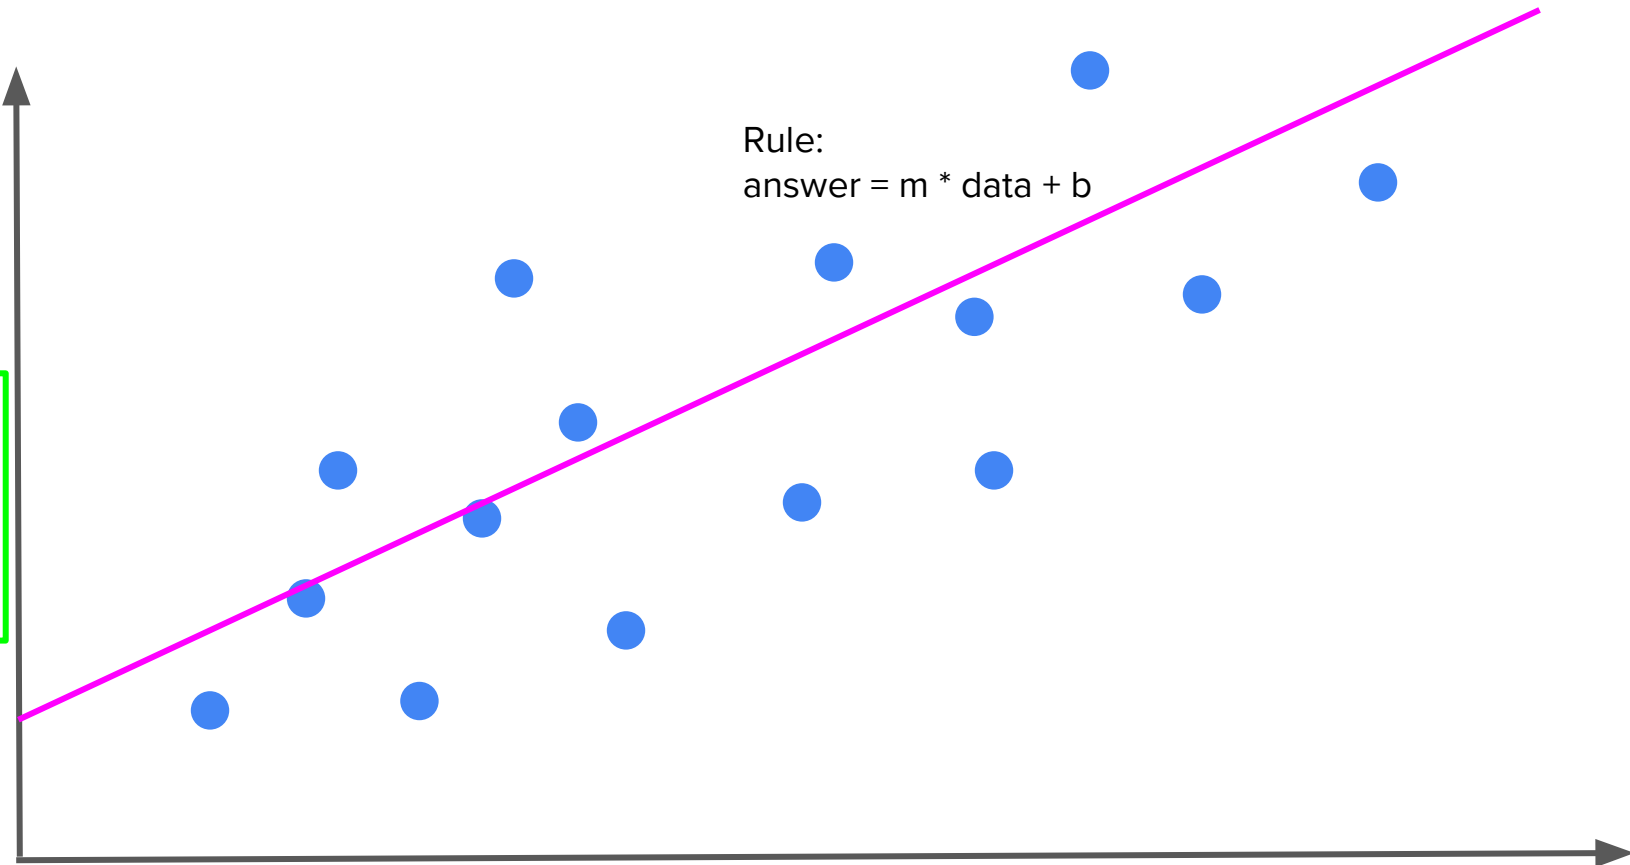
Rule:  
$$\text{answer} = m * \text{data} + b$$

List Price

Square Footage

Data

Answers





Model:

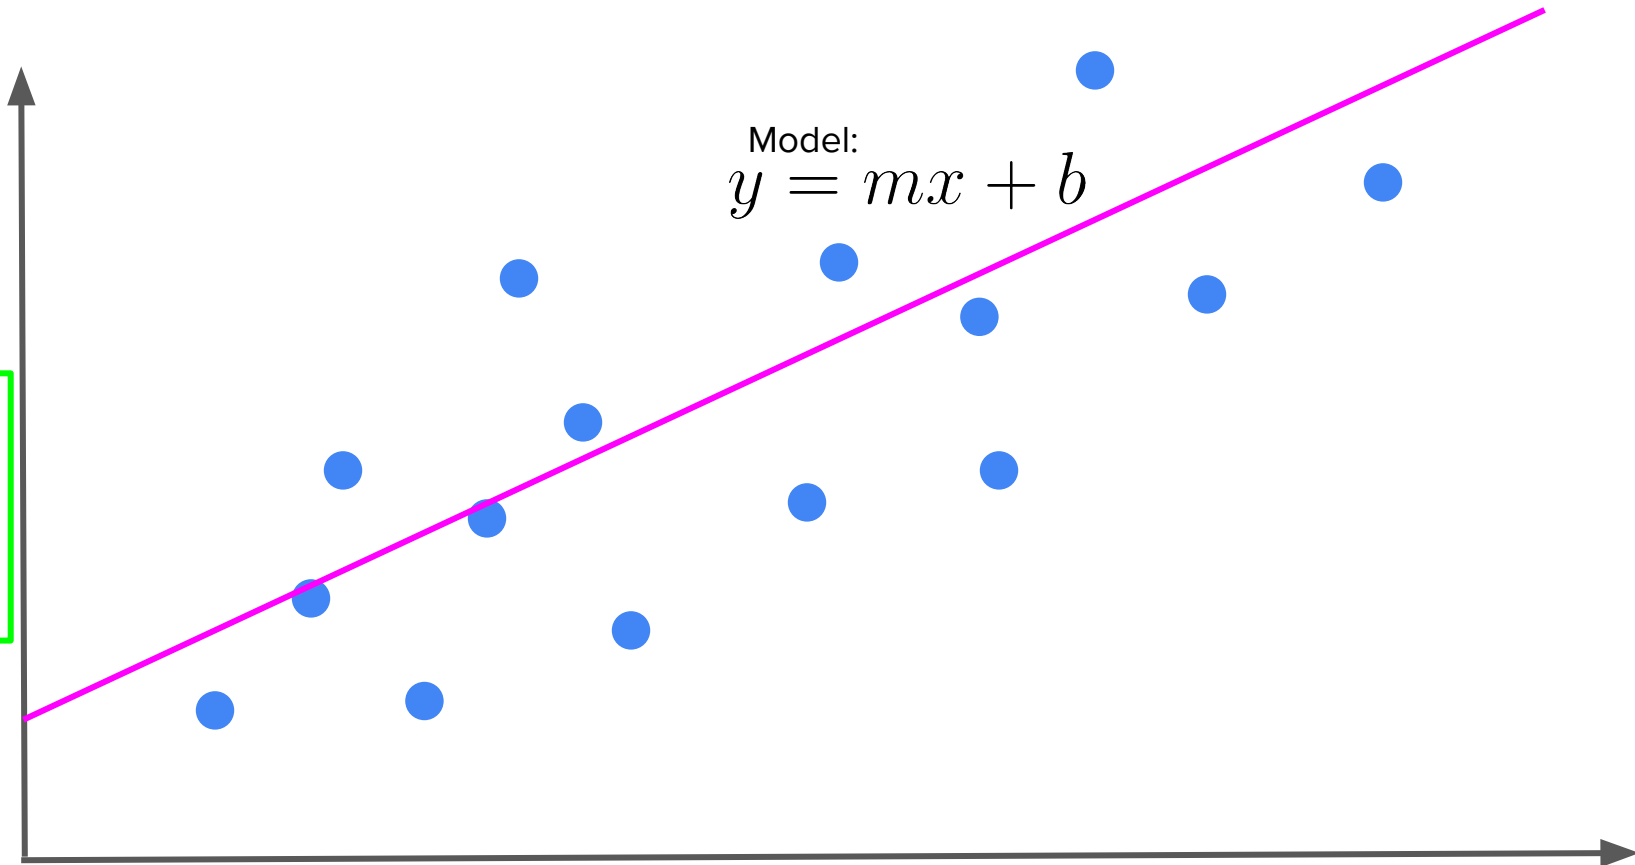
$$y = mx + b$$

List Price

Square Footage

Data

Answers



## Let's Generalize

- Imagine we have  $n$  samples or data points indexed by  $i$ .
- For each sample, we have  $p$  features (known “piece of information” about the sample) indexed by  $j$ .
- A linear model will look like:

$$y_i = \beta_0 + \beta_1 X_{i1} + \beta_2 X_{i2} + \dots + \beta_p X_{ip}$$

Where  $\beta_j$  denotes the  $j$ th model parameter/coefficient.

# Let's Generalize

- Imagine we have  $n$  samples or data points indexed by  $i$ .
- For each sample, we have  $p$  features (known “piece of information” about the sample) indexed by  $j$ .
- A linear model will look like:

$$y_i = \beta_0 + \beta_1 X_{i1} + \beta_2 X_{i2} + \dots + \beta_p X_{ip}$$


Known Answer

Where  $\beta_j$  denotes the  $j$ th model parameter/coefficient.

“Rules”  
aka this is what we want to learn!

# Let's Generalize

- Imagine we have  $n$  samples or data points indexed by  $i$ .
- For each sample, we have  $p$  features (known “piece of information” about the sample) indexed by  $j$ .
- A linear model will look like:

$$y_i = \beta_0 + \beta_1 X_{i1} + \beta_2 X_{i2} + \dots + \beta_p X_{ip}$$

Known Answer

Where  $\beta_j$  denotes the  $j$ th model parameter/coefficient.

“Rules”  
aka this is what we want to learn!

Bias term (y-intercept)

Known Data

Let's Generalize - single sample  $i$

$$y_i = \beta_0 + \beta_1 X_{i1} + \beta_2 X_{i2} + \dots + \beta_p X_{ip}$$

Let's Generalize - single sample  $i$

$$y_i = \beta_0 + \beta_1 X_{i1} + \beta_2 X_{i2} + \dots + \beta_p X_{ip}$$

$$y_i = \sum_{j=0}^p \beta_j X_{ij}$$

where we assume that  $X_{i0} = 1$  (for mathematical convenience).

Let's Generalize -  $n$  samples

$$y_i = \sum_{j=0}^p \beta_j X_{ij}$$

$$\begin{bmatrix} y_1 \\ y_2 \\ \cdot \\ \cdot \\ \cdot \\ y_n \end{bmatrix} = \begin{bmatrix} X_{10} & X_{11} & X_{12} & \dots & X_{1p} \\ X_{20} & X_{21} & X_{22} & \dots & X_{2p} \\ \vdots & \vdots & \ddots & \vdots & \\ X_{n0} & X_{n1} & X_{n2} & \dots & X_{np} \end{bmatrix} \cdot \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \cdot \\ \cdot \\ \beta_p \end{bmatrix}$$

$$\vec{y} = \mathbf{X} \vec{\beta}$$

Let's Generalize -  $n$  samples

$$y_i = \sum_{j=0}^p \beta_j X_{ij}$$

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} X_{10} & X_{11} & X_{12} & \dots & X_{1p} \\ X_{20} & X_{21} & X_{22} & \dots & X_{2p} \\ \vdots & \vdots & \ddots & \vdots & \\ X_{n0} & X_{n1} & X_{n2} & \dots & X_{np} \end{bmatrix} \cdot \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \vdots \\ \beta_p \end{bmatrix}$$

Known Data (features)

$$\vec{y} = \mathbf{X} \vec{\beta}$$

“Rules”  
aka this is what we want to learn!

Known Answers



# The ML Recipe

1. Think up some model
2. Feed **data** into the model and make predictions.
3. Calculate the loss between predictions and true values.
4. Determine the model parameters that produce the minimum loss.

# The ML Recipe

1. Think up some model
2. Feed **data** into the model and make predictions.
3. Calculate the loss between predictions and true values.
4. Determine the model parameters that produce the minimum loss.

Housekeeping:

The model *predictions* are often denoted  $\hat{y}_i$

The *ground truth* / *answers* are often denoted  $y_i$  (no hat).

# The ML Recipe

1. Think up some model
2. Feed **data** into the model and make predictions.

$$\hat{\vec{y}} = \mathbf{X}\vec{\beta}$$

# The ML Recipe

1. Think up some model
2. Feed **data** into the model and make predictions.
3. Calculate the loss between predictions and true values.

Mean Squared Error (MSE):

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n \left( y_i - \vec{X}_i \cdot \vec{\beta} \right)^2$$

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n \left( y_i - \sum_{j=0}^p X_{ij} \beta_j \right)^2$$

# The ML Recipe

1. Think up some model
2. Feed **data** into the model and make predictions.
3. Calculate the loss between predictions and true values.

Mean Squared Error (MSE):

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

true value                      prediction

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n \left( y_i - \vec{X}_i \cdot \vec{\beta} \right)^2$$

data                      Unknown!

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n \left( y_i - \sum_{j=0}^p X_{ij} \beta_j \right)^2$$

# The ML Recipe

1. Think up some model
2. Feed **data** into the model and make predictions.
3. Calculate the loss between predictions and true values.
4. Determine the model parameters that produce the minimum loss.
  - a. Take the derivative of your loss function with respect to the model parameters.
  - b. Set it equal to zero.
  - c. Solve for the model parameters.

# The ML Recipe

1. Think up some model
2. Feed **data** into the model and make predictions.
3. Calculate the loss between predictions and true values.
4. Determine the model parameters that produce the minimum loss.
  - a. Take the derivative of your loss function with respect to the model parameters.
  - b. Set it equal to zero.
  - c. Solve for the model parameters.

$$\frac{\partial \mathcal{L}}{\partial \vec{\beta}} = 0$$

# The ML Recipe

1. Think up some model
2. Feed **data** into the model and make predictions.
3. Calculate the loss between predictions and true values.
4. Determine the model parameters that produce the minimum loss.
  - a. Take the derivative of your loss function with respect to the model parameters.
  - b. Set it equal to zero.
  - c. Solve for the model parameters.

$$\frac{\partial \mathcal{L}}{\partial \vec{\beta}} = 0$$

$$\vec{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \vec{y}$$



# The ML Recipe - Meta Comments

1. Think up some model
  - a. We chose a linear model. There are many others!
2. Feed **data** into the model and make predictions.
  - a. We chose what features to use.
3. Calculate the loss between predictions and true values.
  - a. We chose Mean Square Error as our loss. There are many others!
4. Determine the model parameters that produce the minimum loss.
  - a. We analytically determined the parameters. There are many other ways to determine them!
  - b. Also, sometimes you cannot analytically determine the parameters.


# The ML Recipe - Meta Comments

1. Think up some model
  - a. We chose a linear model. There are many others!
2. Feed **data** into the model and make predictions.
  - a. We chose what features to use.
3. Calculate the loss between predictions and true values.
  - a. We chose Mean Square Error as our loss. There are many others!
4. Determine the model parameters that produce the minimum loss.
  - a. We analytically determined the parameters. There are many other ways to determine them!
  - b. Also, sometimes you cannot analytically determine the parameters.

While we denote step 1 as “the model”, all of these steps involve choices that impact the eventual model that is used for prediction.


# Inference vs. Prediction

- So far, we have focused on being able to build a model that predicts things.



# Inference vs. Prediction

- Inference: inspect your model to learn about your system.



# Inference vs. Prediction

- Inference: inspect your model to learn about your system.
- Examples:
  - Which feature is most important for accurate predictions?
  - If I increase feature Z by 10%, how will that change the prediction?
  - What is the uncertainty in my prediction?
  - What is the uncertainty in a model parameter?

# Inference vs. Prediction

- Inference: inspect your model to learn about your system.
- Examples:
  - Which feature is most important for accurate predictions?
  - If I increase feature Z by 10%, how will that change the prediction?
  - What is the uncertainty in my prediction?
  - What is the uncertainty in a model parameter?
- This is often more difficult and requires statistical guarantees.
- This is well-studied for linear models but not so for more complicated models.
- There is an inherent tradeoff between model complexity and interpretability.
- Model complexity can be correlated with model *performance*, so there can be a tradeoff between performance and interpretability.

# Nonlinear Features for Linear Models

- We can do whatever we want to the features  $\mathbf{X}$ .

$$y_i = \sum_{j=0}^p \beta_j X_{ij}$$

- We can square a feature, we can multiple features by each other, we can apply a sine, etc...

$$y_i = \beta_0 + \beta_1 X_{i1} + \beta_2 X_{i1}^2 + \beta_3 X_{i1} X_{i2} + \beta_4 \sin(X_{i3})$$

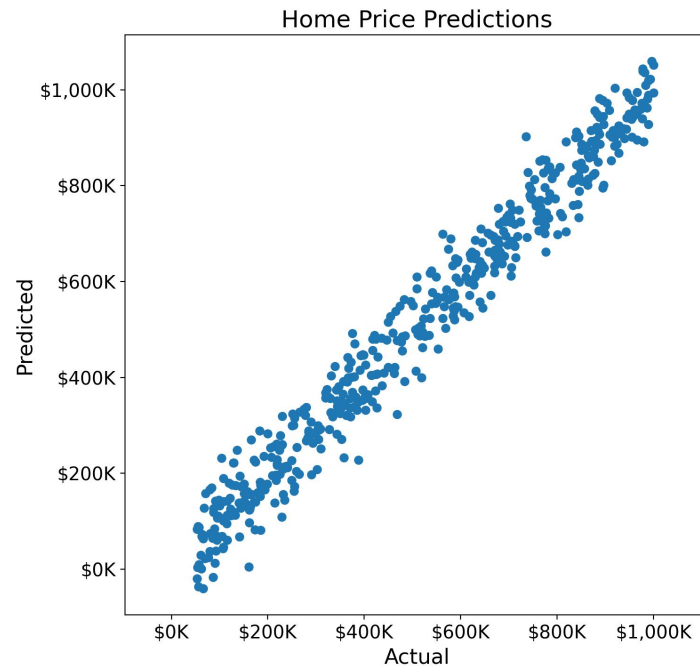
- While these features are nonlinear, the model is linear in the parameters  $\beta$ .

# Evaluating Regression Models

---

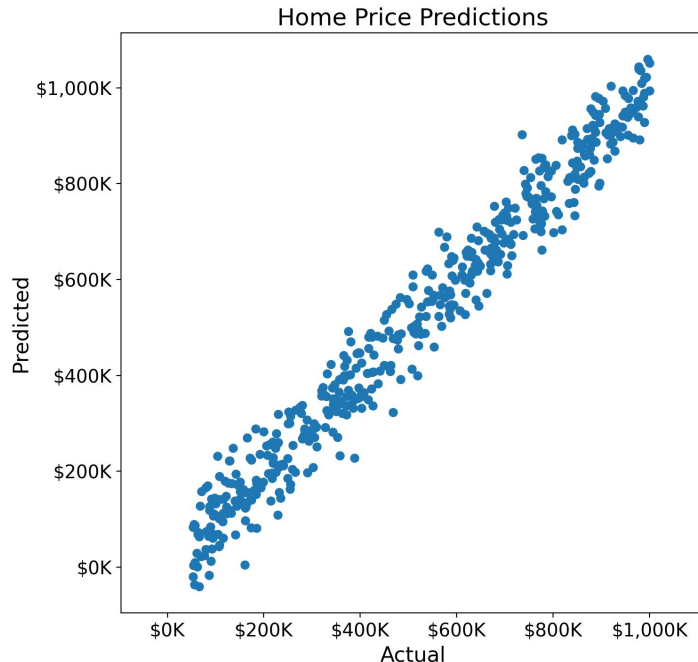


# How do we know how good our prediction is?



# How do we know how good our prediction is?

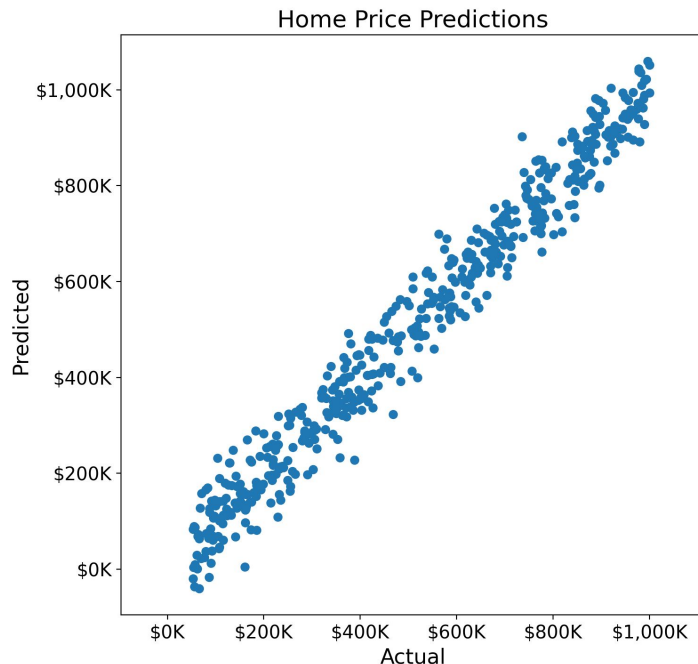
- The model's performance is based on the context in which the model is used.
- We often try to collapse “performance” down to a single number.
- There are many statistical measures of model quality which may or may not be useful in practice.



# How do we know how good our prediction is?

- If we are minimizing the Mean Squared Error (MSE), then we can use MSE as a performance metric.
- Outliers will have a large effect on this metric.

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$



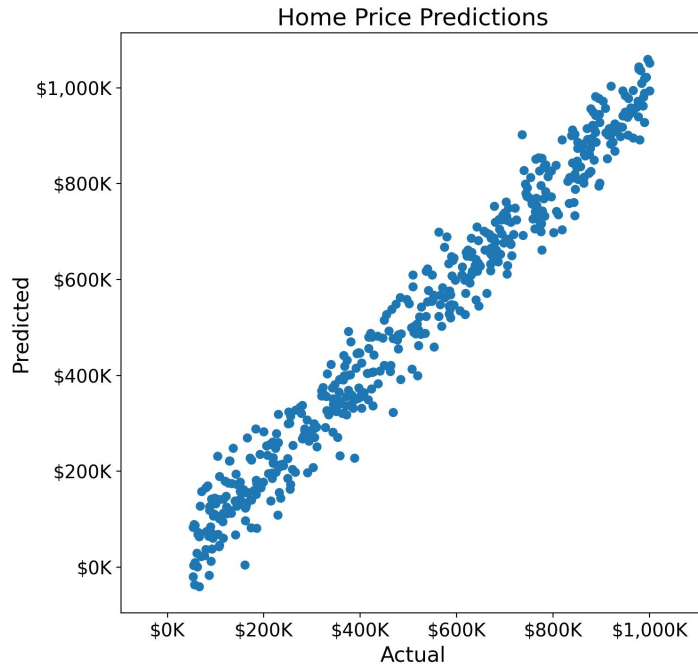
# Regression Metrics

- Mean Squared Error - minimizes outliers, good statistical properties, hard to interpret.

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- Mean Absolute Error - easy to interpret, all loss is treated equally, not robust to outliers.

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$



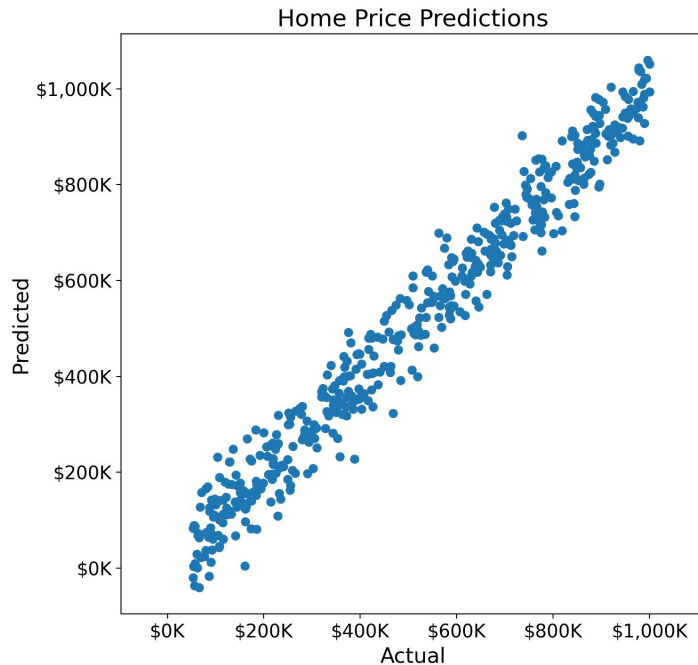
# Regression Metrics

- $\mathcal{R}^2$ / coefficient of determination - statistically interpretable as the proportion of variance that's predictable

$$\mathcal{R}^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

- Mean Absolute Percentage Error - interpretable, error is independent of scale.

$$MAPE = \frac{1}{n} \sum_{i=1}^n \left| \frac{y_i - \hat{y}_i}{y_i} \right|$$



# Scientific Computing

---