

# NYU FRE 7773 - Week 6

---

*Machine Learning in Financial Engineering*

Ethan Rosenthal

# Time Series Machine Learning


---

*Machine Learning in Financial Engineering*

Ethan Rosenthal


Time

---



# Uber trips


<https://eng.uber.com/neural-networks/>



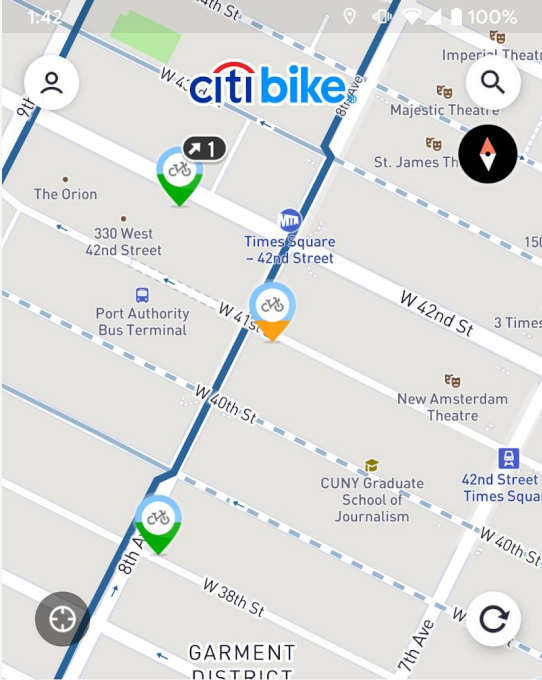
# Amazon weekly item sales

## DeepAR: Probabilistic Forecasting with Autoregressive Recurrent Networks

<https://arxiv.org/abs/1704.04110>



No Citi Bikes



0 seconds ago

W 41 St & 8 Ave

2

Bikes

55

Docks

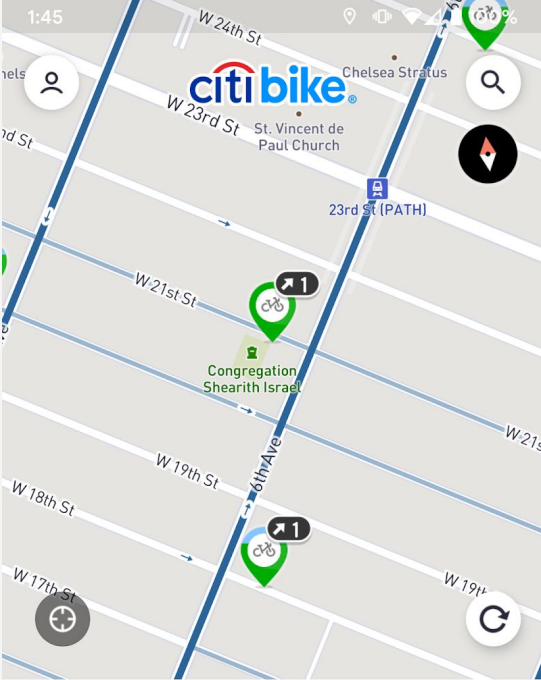


Unlock a Bike



Plan a Ride

No Docks



13 seconds ago

W 21 St & 6 Ave

49

Bikes

0


Docks




Unlock a Bike



Plan a Ride








# “Classical” Time Series Modeling

---



*State Space Models*



*Nonlinear PDE's*



*ARIMA*



*Gauss-Markov Theorem*




$$\begin{array}{ccccc}
 (1 - \phi_1 B - \dots - \phi_p B^p) & (1 - B)^d y_t & = & c + (1 + \theta_1 B + \dots + \theta_q B^q) \varepsilon_t \\
 \uparrow & \uparrow & & \uparrow \\
 \text{AR}(p) & d \text{ differences} & & \text{MA}(q)
 \end{array}$$




**Can you skip  
all of this?**

# Inference



# Prediction



# Where's the X Matrix?

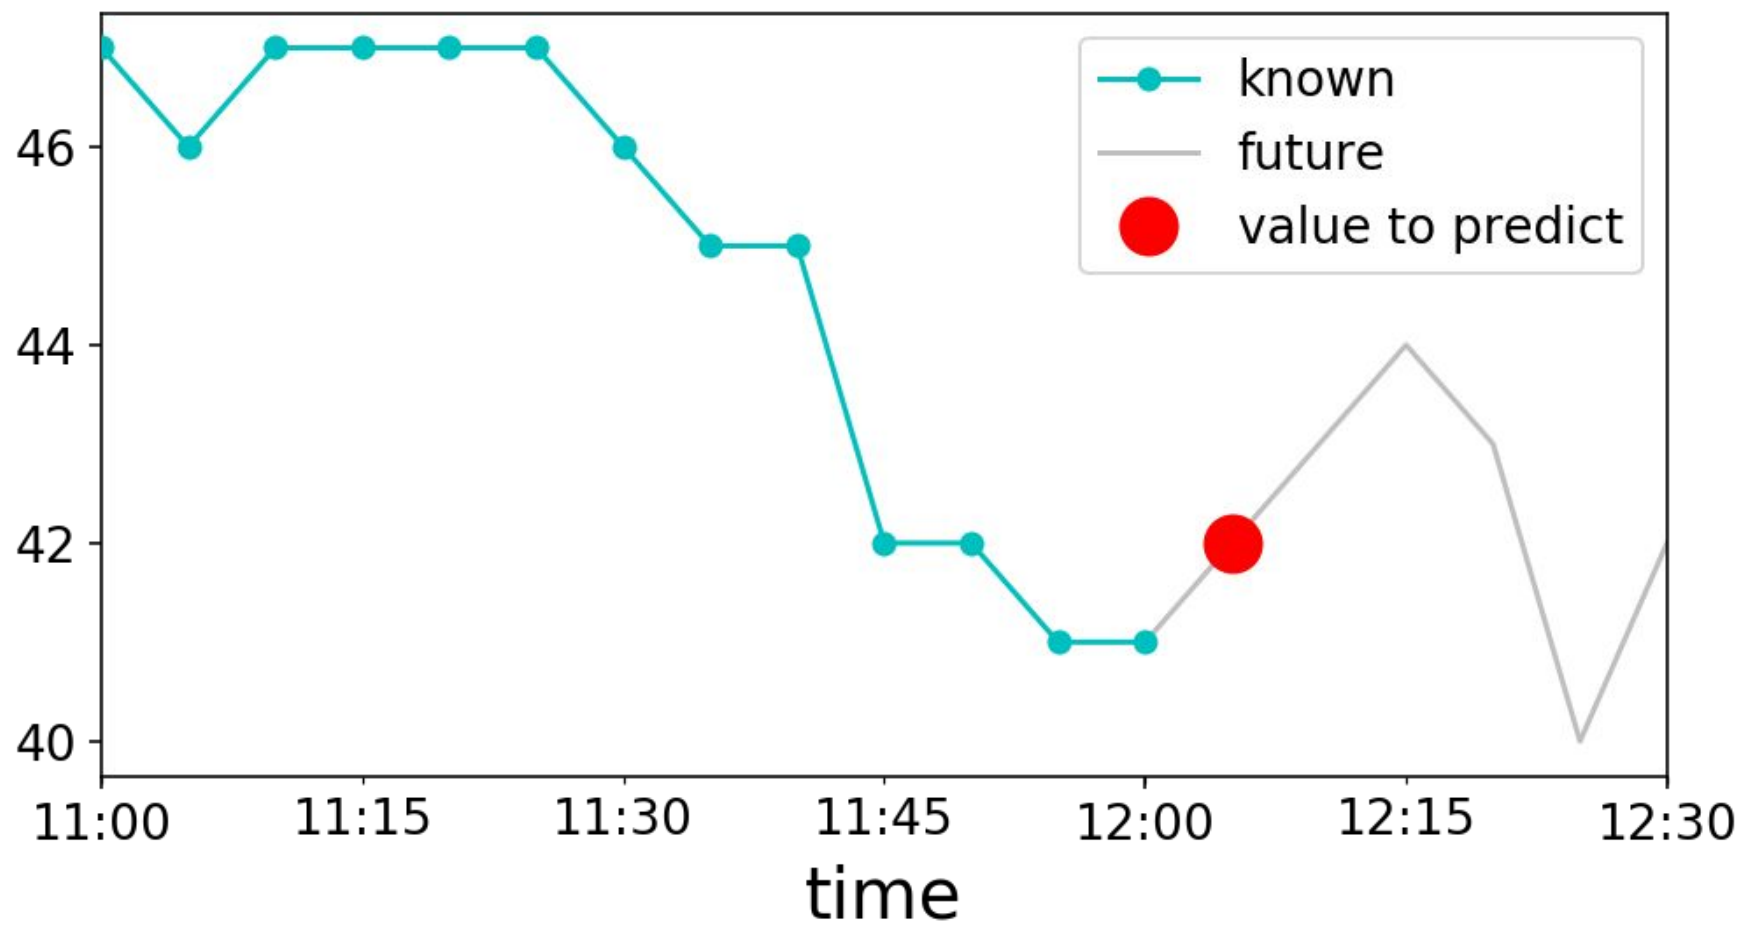
---

```
f.fit(X, y)  
y_new = f.predict(X_new)
```

In the beginning, there was ***y***








```
X = []  
for idx in range(len(y) - window):  
    X.append(y[idx:idx + window])  
  
X = np.array(X)
```

```
y: array([0,  
         1,  
         2,  
         3,  
         4,  
         5,  
         6,  
         7,  
         8,  
         9,  
        10])
```

```
X: array([[0, 1, 2, 3, 4],  
         [1, 2, 3, 4, 5],  
         [2, 3, 4, 5, 6],  
         [3, 4, 5, 6, 7],  
         [4, 5, 6, 7, 8],  
         [5, 6, 7, 8, 9]])
```



$$\mathbf{X}\beta = \hat{\mathbf{y}}$$

$$\begin{bmatrix} y_0 & y_1 & y_2 & \cdots & y_{w-1} \\ y_1 & y_2 & y_3 & \cdots & y_w \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ y_{t-2-w} & y_{t-1-w} & y_{t-w} & \cdots & y_{t-2} \\ y_{t-1-w} & y_{t-w} & y_{t-w+1} & \cdots & y_{t-1} \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_{w-2} \\ \beta_{w-1} \end{bmatrix} = \begin{bmatrix} \hat{y}_w \\ \hat{y}_{w+1} \\ \vdots \\ \hat{y}_{t-1} \\ \hat{y}_t \end{bmatrix}$$

# Modeling

- Not limited to Linear Regression. Use trees, neural nets, whatever you want!
- Not limited to regression. Classification, quantile regression, etc...

# Preprocessing and Feature Engineering

- Differencing
- Rolling Mean
- Filters
- Fourier components
- Seasonal lags
- etc...

## Adding “Exogenous” Features

```
X = []  
for idx in range(len(y) - window):  
    X.append(y[idx:idx + window])
```

```
X = np.array(X)  
X = np.hstack((X, X_features))
```

**X:** array([[0, 1, 2, 3, 4, 0.5, -0.1],  
[1, 2, 3, 4, 5, 2.3, 0.2],  
[2, 3, 4, 5, 6, -0.2, 0.4],  
[3, 4, 5, 6, 7, 0.9, 1.1],  
[4, 5, 6, 7, 8, 1.2, 0.5],  
[5, 6, 7, 8, 9, -0.7, -0.2]])

**y:** array([0,  
1,  
2,  
3,  
4,  
5,  
6,  
7,  
8,  
9,  
10])

**Lag Features**      **Extra Features**

The diagram illustrates the construction of the feature matrix X from the target variable y. The first 5 columns of X are labeled 'Lag Features' and represent the sequence of y values from index idx to idx+5. The last 2 columns are labeled 'Extra Features' and represent additional information. Arrows indicate that each row of X corresponds to a specific index in y, with the lag features being the previous 5 values of y and the extra features being additional data points.

```
X = []  
for idx in range(len(y) - window):  
    X.append(y[idx:idx + window])
```

```
X = np.array(X)  
X = np.hstack((X, X_features))
```

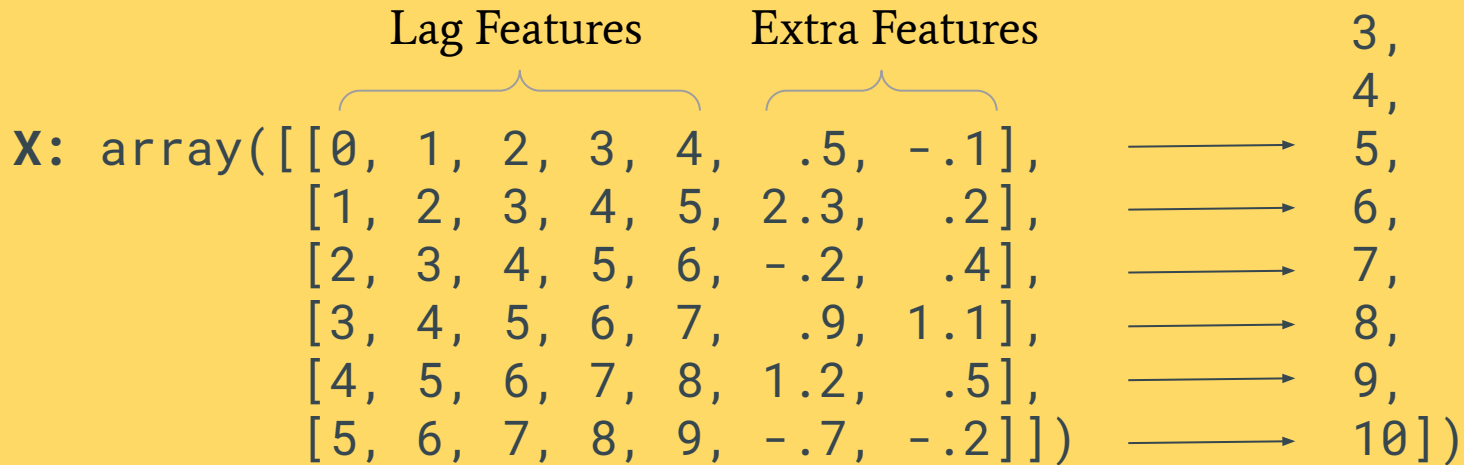
Adding “Exogenous” Features

Complicated to  
construct!

**X:** array([[0, 1, 2, 3, 4, .5, -.1],  
[1, 2, 3, 4, 5, 2.3, .2],  
[2, 3, 4, 5, 6, -.2, .4],  
[3, 4, 5, 6, 7, .9, 1.1],  
[4, 5, 6, 7, 8, 1.2, .5],  
[5, 6, 7, 8, 9, -.7, -.2]])

**Lag Features**      **Extra Features**

**y:** array([0,  
1,  
2,  
3,  
4,  
5,  
6,  
7,  
8,  
9,  
10])





# Recap

- Take your time series  $\mathbf{y}$ .
- Treat each point in  $\mathbf{y}$  as a point that you want to predict.
- Construct  $\mathbf{X}$  from any data you want that comes *prior* to the point in  $\mathbf{y}$  that you want to predict.

$$\mathbf{X}_t = \mathbf{y}_{t' < t}$$

- Fit a regression model on  $\mathbf{X}$  and  $\mathbf{y}$ .
- For each point in  $\mathbf{y}$ , use model to predict the next point.


$$\hat{y}_t = f(\mathbf{X}_t)$$

$$\hat{y}_t = f(\mathbf{y}_{t' < t})$$


# Forecasting




# Recursive Forecasting




# Recursive Forecasting



# Recursive Forecasting



# Recursive Forecasting




```
def recursive_forecast(model, input_data, num_points_in_future):  
    for point in range(num_points_in_future):  
        prediction = model.predict(input_data)  
        # Append prediction to the input data  
        input_data = np.hstack((input_data, prediction))  
  
    return prediction
```

Optimize for  
next step

Pray recursive  
steps work


# Horizon Forecasting



More info: “Machine learning strategies for multi-step-ahead time series forecasting”,  
Souhaib Ben Taieb [https://souhaib-bentaieb.com/papers/2014\\_phd.pdf](https://souhaib-bentaieb.com/papers/2014_phd.pdf)



# Horizon Forecasting



```
y: array([ 0.,  
          1.,  
          2.,  
          3.,  
          4.,  
          5.,  
          6.,  
          7.,  
          8.,  
          9.,  
         10.])  
y_horizon: array([[ 0.,  1.,  2.],  
                  [ 1.,  2.,  3.],  
                  [ 2.,  3.,  4.],  
                  [ 3.,  4.,  5.],  
                  [ 4.,  5.,  6.],  
                  [ 5.,  6.,  7.],  
                  [ 6.,  7.,  8.],  
                  [ 7.,  8.,  9.],  
                  [ 8.,  9., 10.],  
                  [nan, nan, nan],  
                  [nan, nan, nan]])
```

# Predicting multiple targets


- [sklearn.multioutput.MultiOutputRegressor](#)
- Train an individual model for each target.
- Pros:
  - Very simple
  - Works with any model
- Cons
  - Resource intensive
  - No sharing of knowledge


# Predicting multiple targets

- Deep learning model with multiple outputs.
- Pros:
  - Potentially less resource intensive
  - Direct optimization
  - Sharing of knowledge
- Cons:
  - All the caveats of deep learning
  - Limited to the horizon


# Evaluation

---






Baseline Model: Future = Past







# Forecasting Views




# Forecasting Views




# Forecasting Views



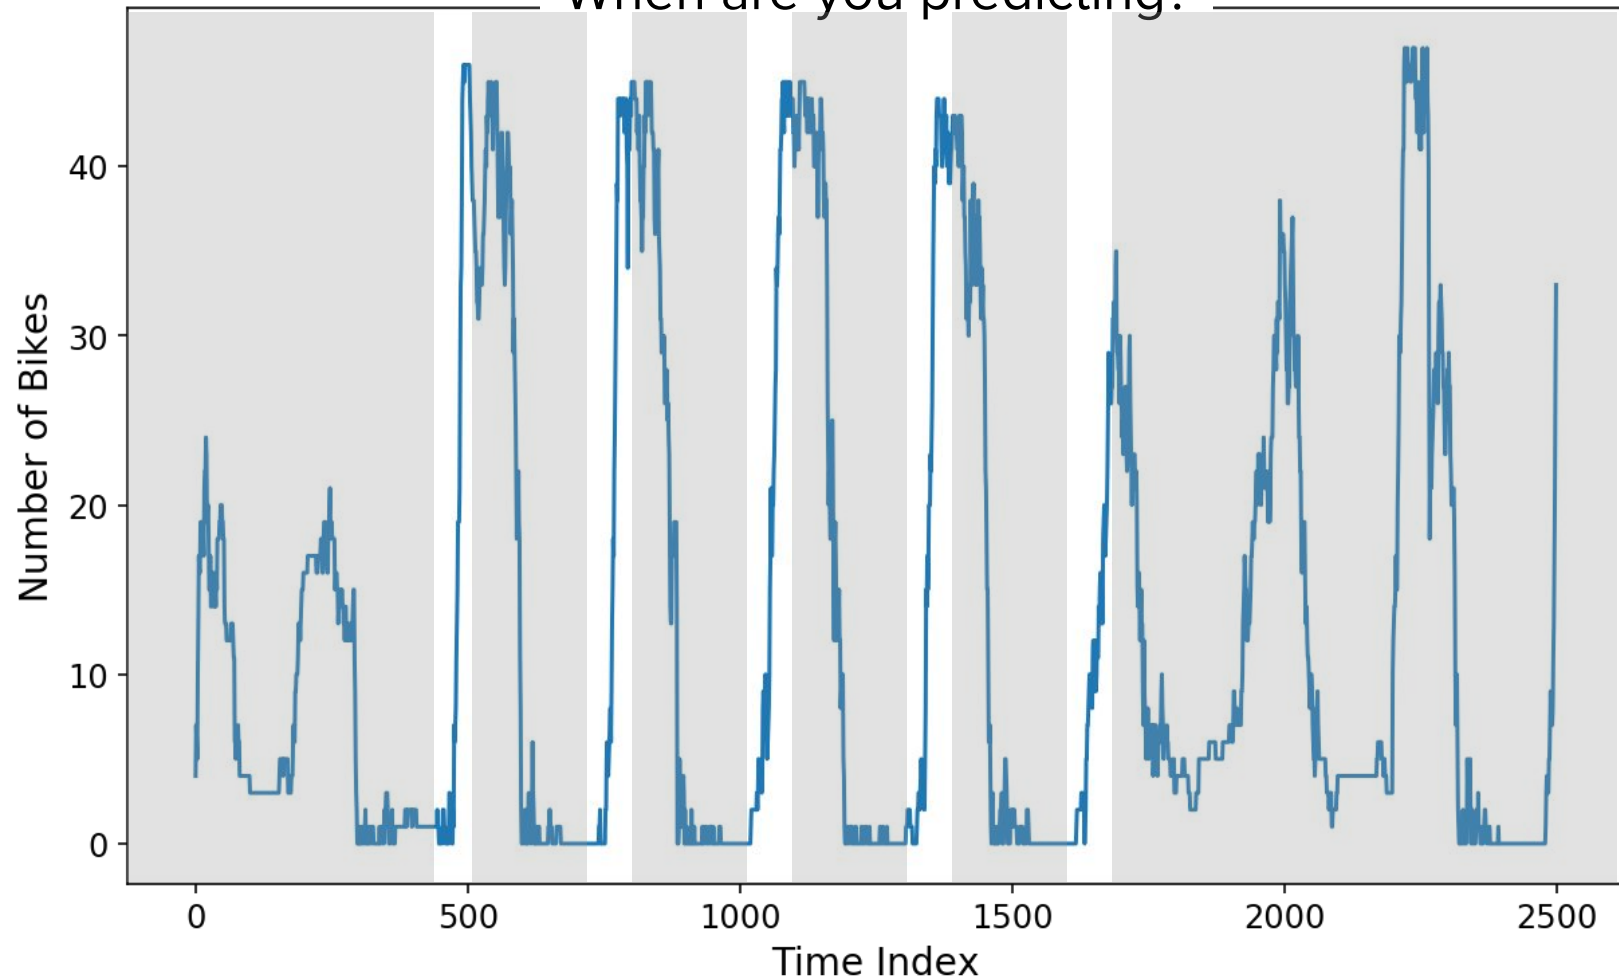
# Metrics Aggregation




When are you predicting?




When are you predicting?



## Regression or Classification?




What do you know when  
you predict?

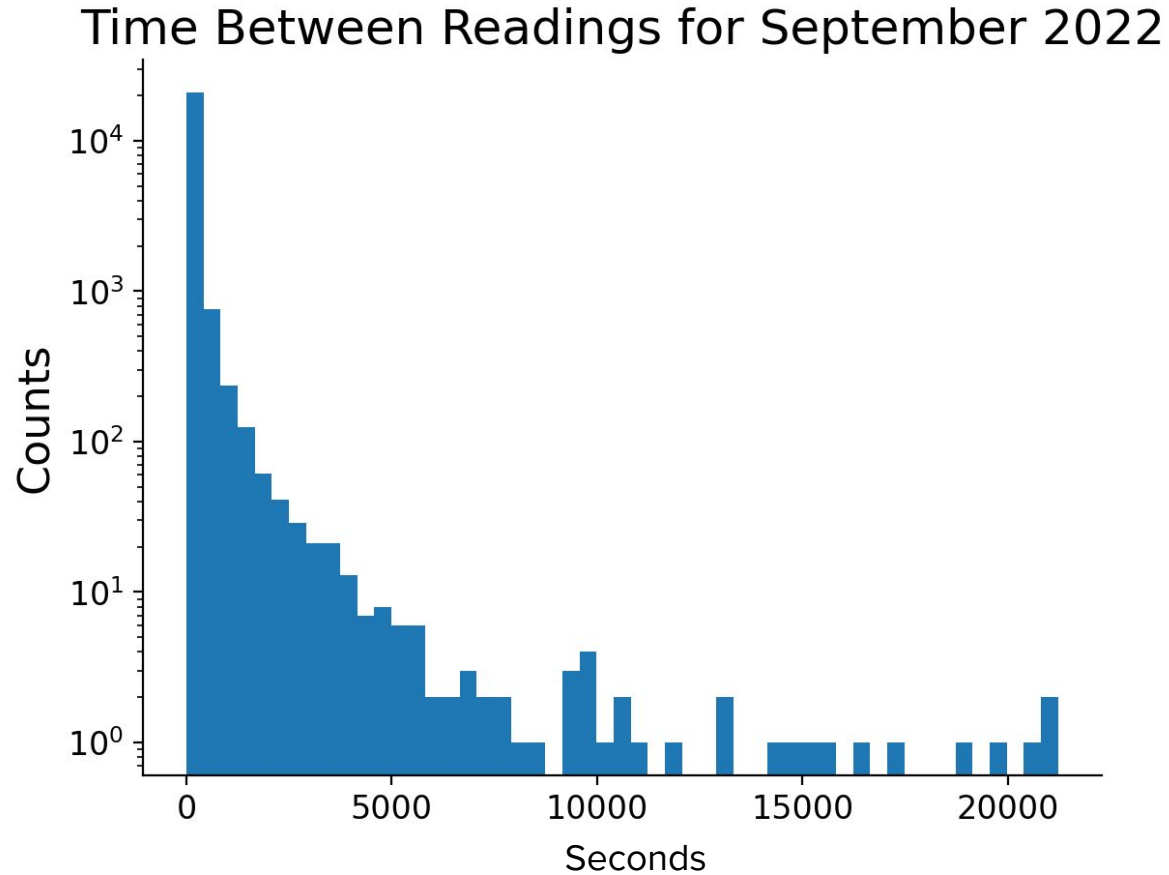




What do you know when  
you predict?




What do you know when  
you predict?




2023

What do you know when  
you predict?



2023

What do you know when  
you predict?



2023

What do you know when  
you predict?

Time Between Readings for September 2023

