

NYU FRE 7773 - Week 12

Machine Learning in Financial Engineering


Jacopo Tagliabue

Introduction to word embeddings

Remember the first lecture?

- **Large language models** are based (sort of) on two simple ideas:

Words are dense vectors



Neural network are good with sequences

The Annotated Transformer

Apr 3, 2018

There is now a [new version](#) of this blog post updated for modern PyTorch.

```
from IPython.display import Image
Image(filename='images/aiayn.png')
```

Attention Is All You Need


Ashish Vaswani* Google Brain avaswani@google.com	Noam Shazeer* Google Brain noam@google.com	Niki Parmar* Google Research nikip@google.com	Jakob Uszkoreit* Google Research usz@google.com
Llion Jones* Google Research llion@google.com	Aidan N. Gomez* [†] University of Toronto aidan@cs.toronto.edu	Lukasz Kaiser* Google Brain lukaszkaiser@google.com	

Remember the first lecture?

- **So far:**

- We saw sparse sequences (Markov models) and sparse vectors (vectorizers!).
- We now introduce **dense vectors**.

Words are dense vectors



Neural network are good with sequences

The Annotated Transformer

Apr 3, 2018

There is now a [new version](#) of this blog post updated for modern PyTorch.

```
from IPython.display import Image
Image(filename='images/aiayn.png')
```

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez*[†]
University of Toronto
aidan@cs.toronto.edu

Lukasz Kaiser*
Google Brain
lukaszkaiser@google.com

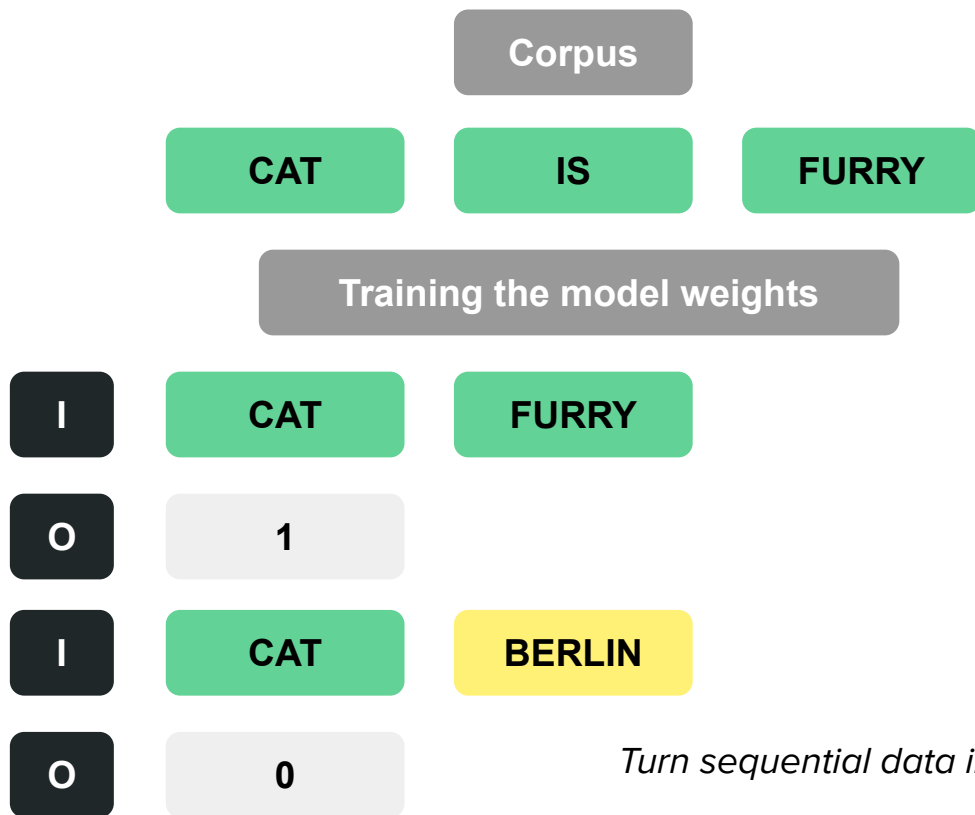
The fundamental intuition of embeddings!

- (Philosophical) **distributional hypothesis**: “words that appear in similar contexts have similar meanings”
- **Computational hypothesis**: learn a classifier that given a word (e.g. cat) tells me how likely is that *another* word appear next to it (Germany, furry) - we don't care about the classifier: we care about the fact that for the classifier to even *succeed*, the cat-vector should be close to the furry-vector!

Word Embeddings
Past, Present and Future

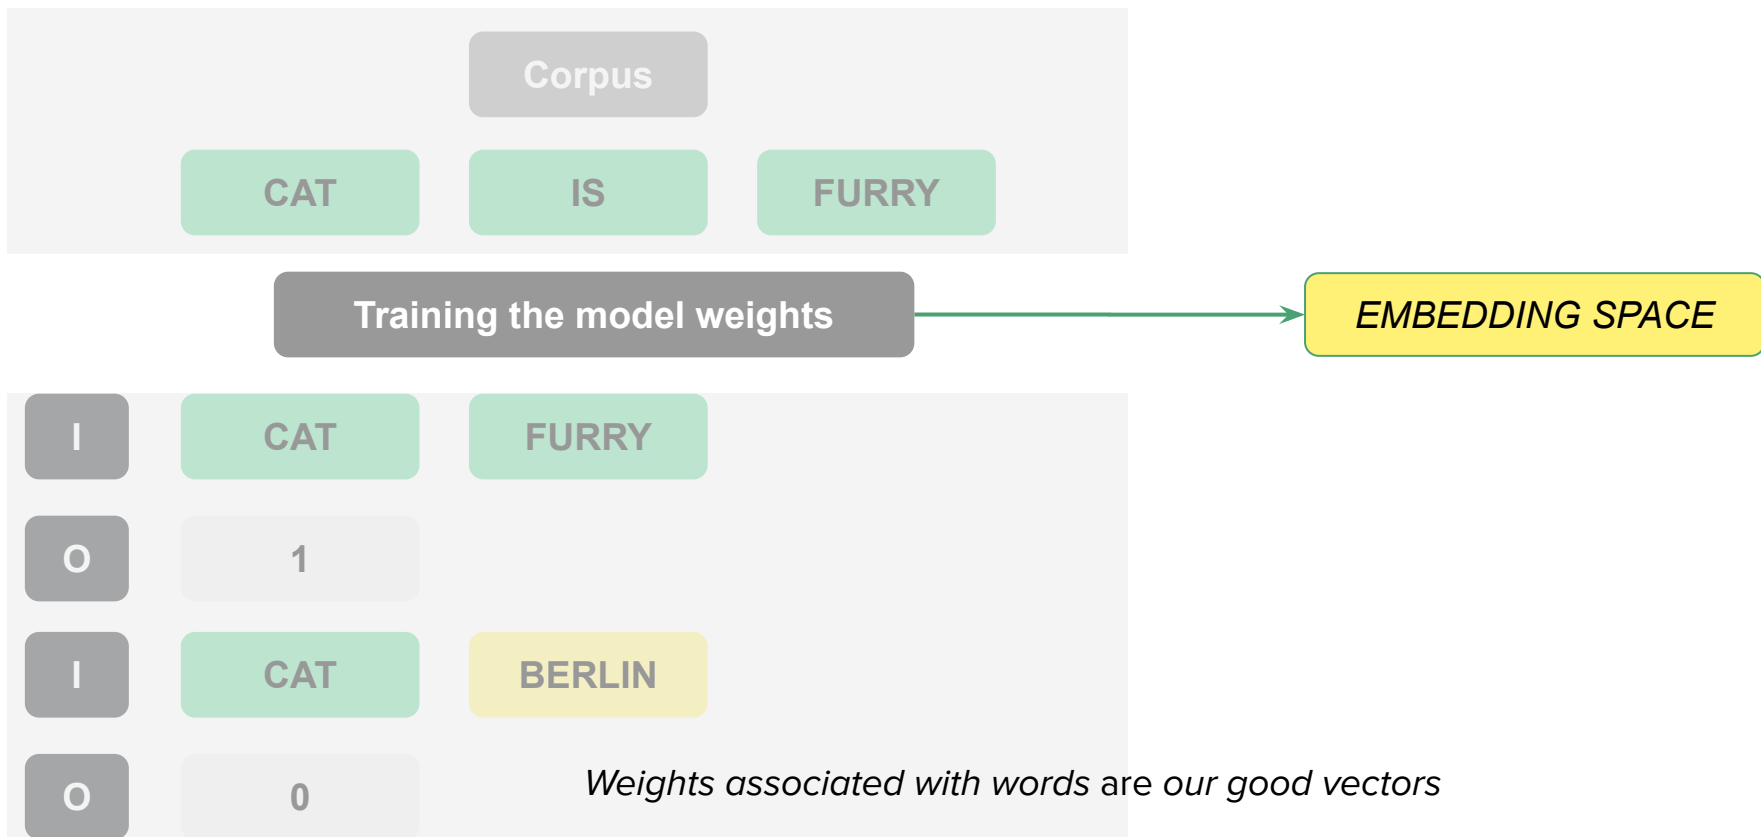


A recipe for learning word embeddings (“word2vec”)

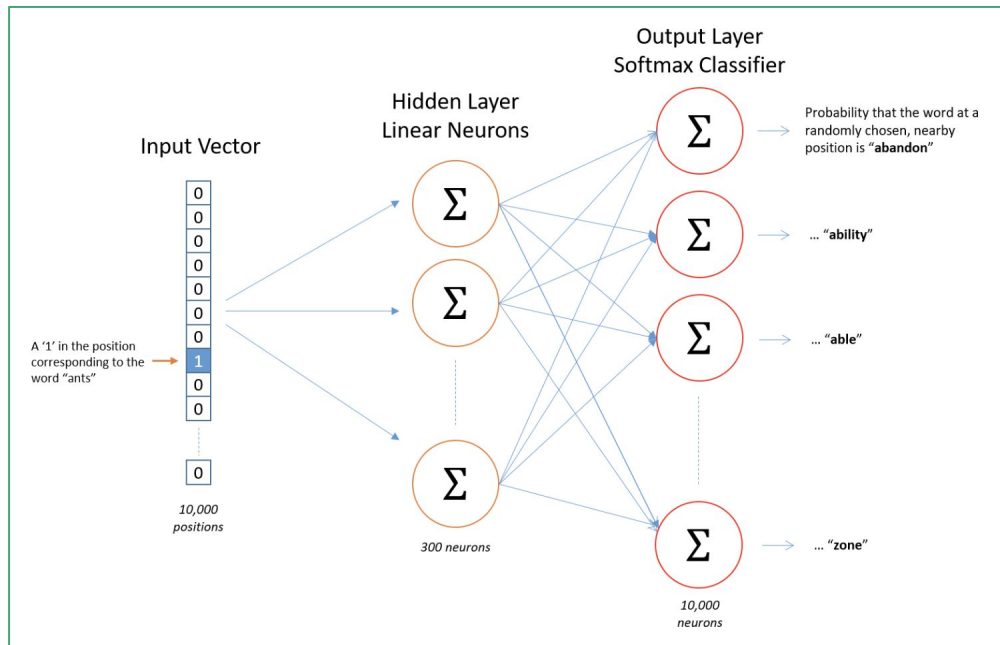


Turn sequential data into a prediction problem

A recipe for learning word embeddings (“word2vec”)



A recipe for learning word embeddings (“word2vec”)



Vec Explained

Based on my previous post: [Vector Representations of Words](#).

Ever wondered what is the opposite of Canada? or if the result of **king-man+woman** is chaos? While first impression tells this post could be filled with some pretty poor jokes,

Word vectors in a *prediction* task

- CORPUS: “The furry cat is on the mat”
- WINDOW LENGTH: 2
- TARGET: “cat”
- INPUT PREPARATION, positive and negative samples

Target	Context	Label
cat	furry	1
cat	the	1
cat	is	1
cat	on	1

Target	Context	Label
cat	Berlin	0
cat	Jacopo	0
cat	ciao	0
cat	table	0

Word vectors in a *prediction* task

- CORPUS: “The furry cat is on the mat”
- WINDOW LENGTH: 2
- TARGET: “cat”
- INPUT PREPARATION, positive and negative samples ($\alpha=0.75$)

$$P_{\alpha}(w) = \frac{\text{count}(w)^{\alpha}}{\sum_{w'} \text{count}(w')^{\alpha}}$$

Target	Context	Label
cat	furry	1
cat	the	1
cat	is	1
cat	on	1

Target	Context	Label
cat	Berlin	0
cat	Jacopo	0
cat	ciao	0
cat	table	0

Word vectors in a *prediction* task

- We have turned a word prediction problem into a binary classification problem
 - Is the context word likely to appear next to the target word?
- Let's define our learning objective:
 - We want to maximize the similarity of (t,c) drawn from the positive examples
 - We want to minimize the similarity of (t,c) drawn from the negative examples

$$L(\theta) = \sum_{(t,c) \in +} \log P(+|t,c) + \sum_{(t,c) \in -} \log P(-|t,c)$$

CHAPTER

6

Vector Semantics and Embeddings

荃者所以在鱼，得鱼而忘荃 Nets are for fish;
Once you get the fish, you can forget the net.
言者所以在意，得意而忘言 Words are for meaning;
Once you get the meaning, you can forget the words
庄子(Zhuangzi), Chapter 26

The parable that Lao Anshu is famous for comes mainly on its forearm. But

Word vectors in a *prediction* task

- Let's define our learning objective:
 - We want to maximize the similarity of (t,c) drawn from the positive examples
 - We want to minimize the similarity of (t,c) drawn from the negative examples

The diagram illustrates the components of the loss function. A blue box labeled "Dot product" has an arrow pointing to the $c \cdot t$ term in the first equation. A green box labeled "Sigmoid" has an arrow pointing to the σ function in the same term. The equations are as follows:

$$= \log \sigma(c \cdot t) + \sum_{i=1}^k \log \sigma(-n_i \cdot t)$$
$$= \log \frac{1}{1 + e^{-c \cdot t}} + \sum_{i=1}^k \log \frac{1}{1 + e^{n_i \cdot t}}$$

Word vectors in a *prediction* task

- **Remember:** we maximize the dot product of the word with the context words, and minimize the dot products of the word with the negative sampled words!
- Training procedure:
 - Random initialization of vectors for N words in the vocabulary.
 - At each step, move embeddings of related words closer in the vector space, and push others further away (using gradient descent).

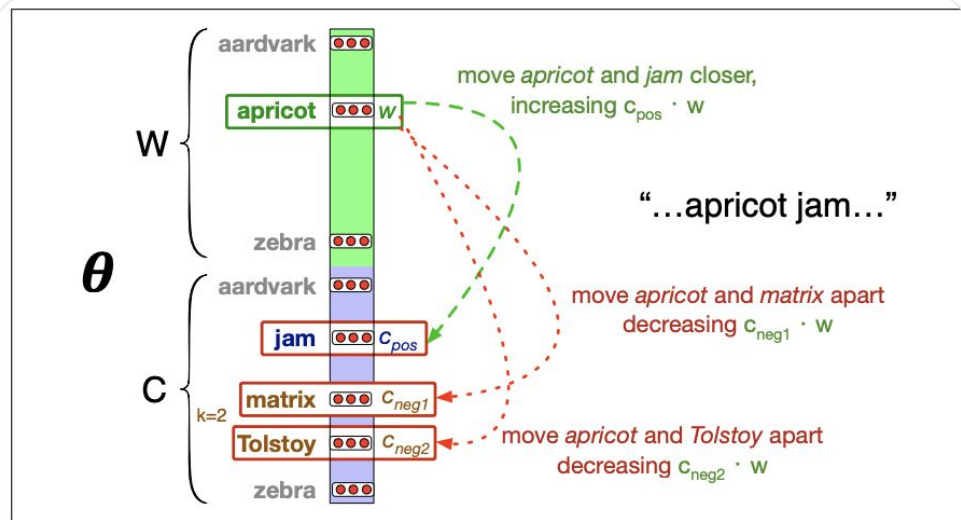



Figure 6.14 Intuition of one step of gradient descent. The skip-gram model tries to shift embeddings so the target embeddings (here for *apricot*) are closer to (have a higher dot product with) context embeddings for nearby words (here *jam*) and further from (lower dot product with) context embeddings for noise words that don't occur nearby (here *Tolstoy* and *matrix*).

Is this a “good” space?

- word2vec tends to capture well similarity between words and some analogical relations - **without any human labels / intervention!**
- Once you have a well-trained embedding space, the offsets between vector embeddings can be used to solve analogies such as: “man : king = women : ?” (*queen*)
 - This is possible since the result of $\text{vector}(\text{'king'}) - \text{vector}(\text{'man'}) + \text{vector}(\text{'woman'})$ is a vector close to $\text{vector}(\text{'queen'})$.

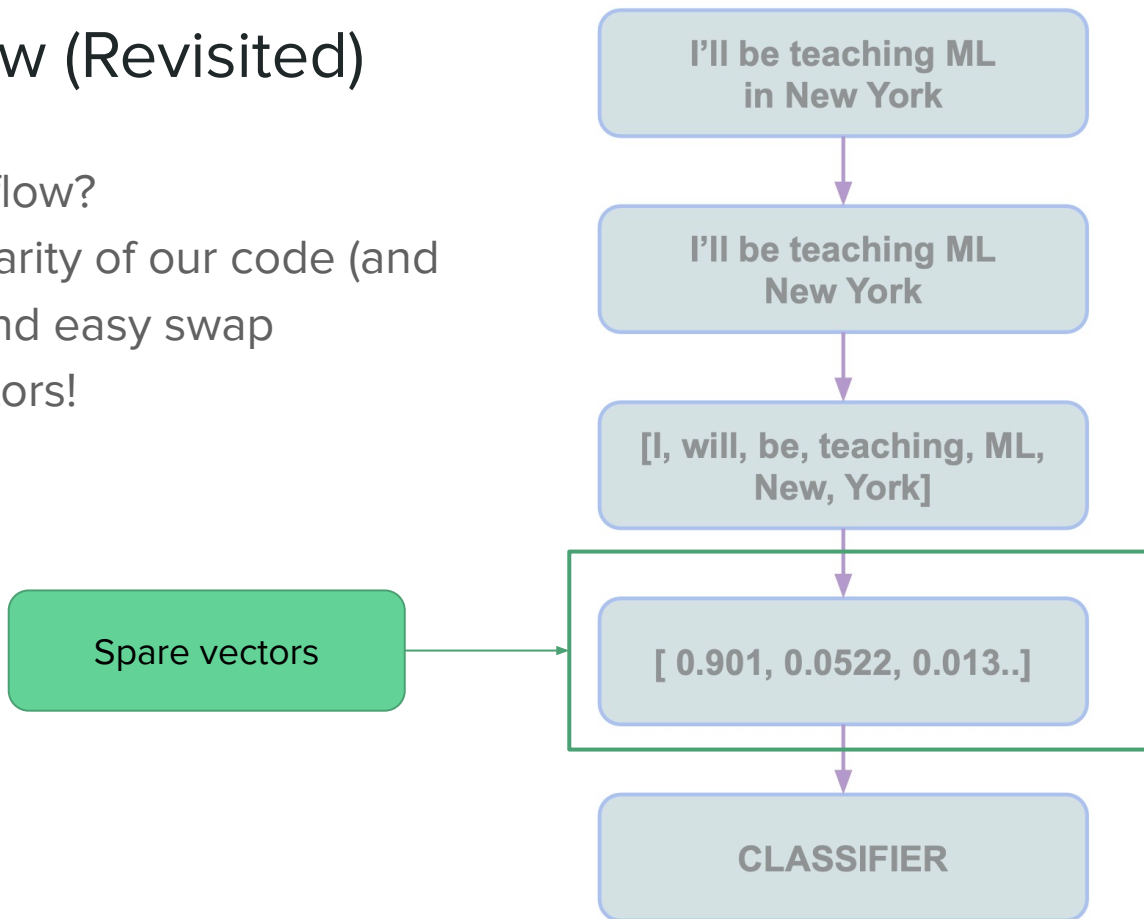


Word2Vec FTW!

- Some pretty cool features of word2vec:
 - Vectors are good for downstream models: if we use 100-dim embeddings as features, a classifier can just learn 100 weights to represent a function.
 - Learning is “self-supervised”, as any text we can find on the internet can be turned into positive/negative pairs without manual intervention.
 - The vector space encodes automatically similarities and analogies.
 - Vectors are “portable”: they can be learned on Wikipedia and applied to finance news.
- The window size influences which “similarity” we care about:
 - Shorter windows leads to representations that are syntactically and semantically similar: “cat” will be very similar to “dog” for example - same topic (pet), same part of speech (noun).
 - Long context windows leads to representations that are topically related but not necessarily syntactically equivalent, such as “cat” and “furry” for example.

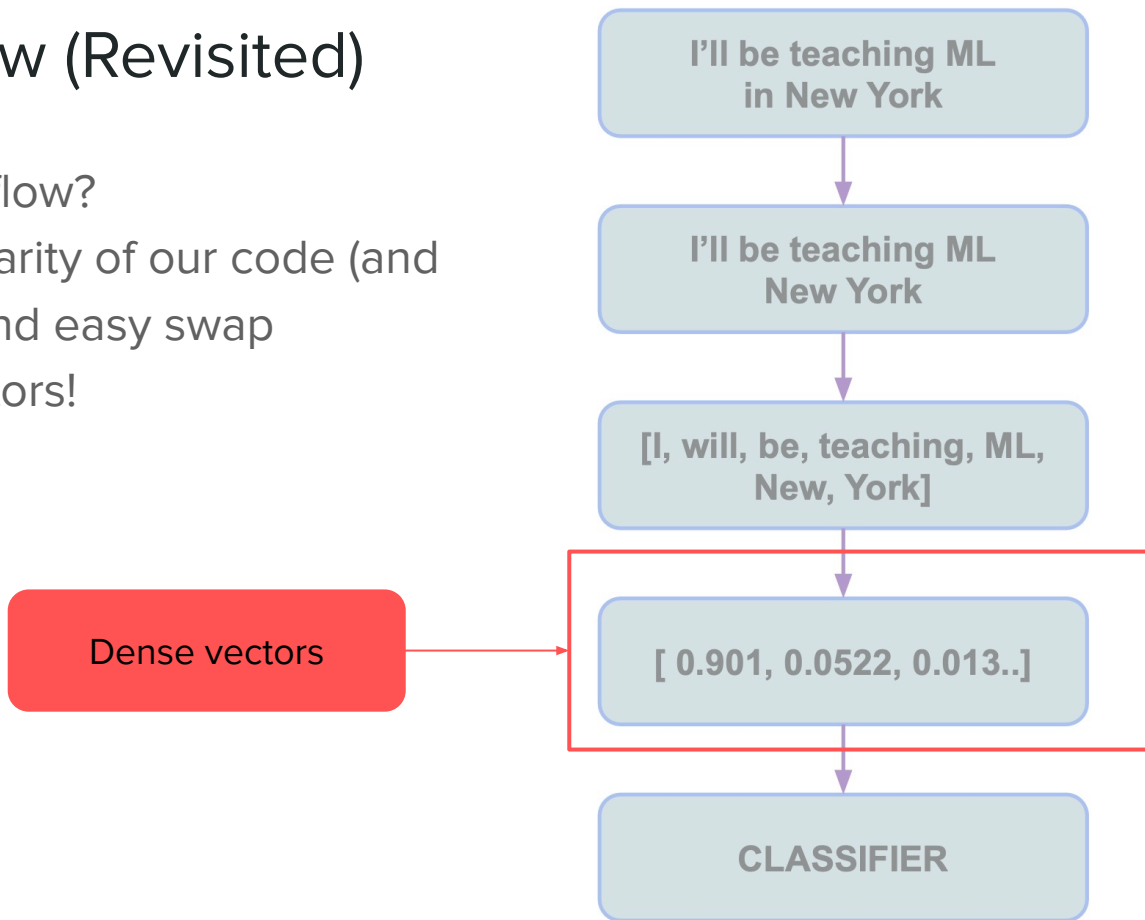
Text Classification Flow (Revisited)

- Remember our NLP workflow?
- We can exploit the modularity of our code (and our Metaflow pipelines) and easily swap “in-and-out” different vectors!




Text Classification Flow (Revisited)

- Remember our NLP workflow?
- We can exploit the modularity of our code (and our Metaflow pipelines) and easily swap “in-and-out” different vectors!



Text Classification Flow (Revisited)

- Remember our NLP workflow?
- We can exploit the modularity of our code (and our Metaflow pipelines) and easily swap “in-and-out” different vectors!




BONUS: thanks to deep architectures, we can actually drastically simplify the first part too!


Everything2vec

Bonus: every time we have a sequence, we have x2vec!


Remember: the distributional hypothesis can be applied whenever we have meaningful sequences of target items (e.g. song playlist, shopping sessions, molecules etc.) - see [MLSys 2022 for a song recSys example!](#)



CAT



IS



FURRY

Song2Vec



Book2Vec



Coding time!