

# NYU FRE 7773 - Week 13

---

*Machine Learning in Financial Engineering*

Jacopo Tagliabue

# Serving predictions

---

*Machine Learning in Financial Engineering*  
Jacopo Tagliabue

# Welcome to the jungle

If your work needs to have an impact, it needs to **RUN OUTSIDE YOUR LAPTOP:**

1. Your code can be **inspected, modified, understood** by others, typically your technical colleagues: you need to write clean, modular, testable code and make your pipeline fully reproducible.
2. Your model can be **trusted** by others, typically, other stakeholders, who may or may not be technical folks: you need to “make sure” the model behaved as designed before pushing it in front of end-users.
3. Predictions can be **consumed** by others, typically anybody with an internet connection: you need to expose your model as an endpoint which returns predictions when supplied with the appropriate parameters.

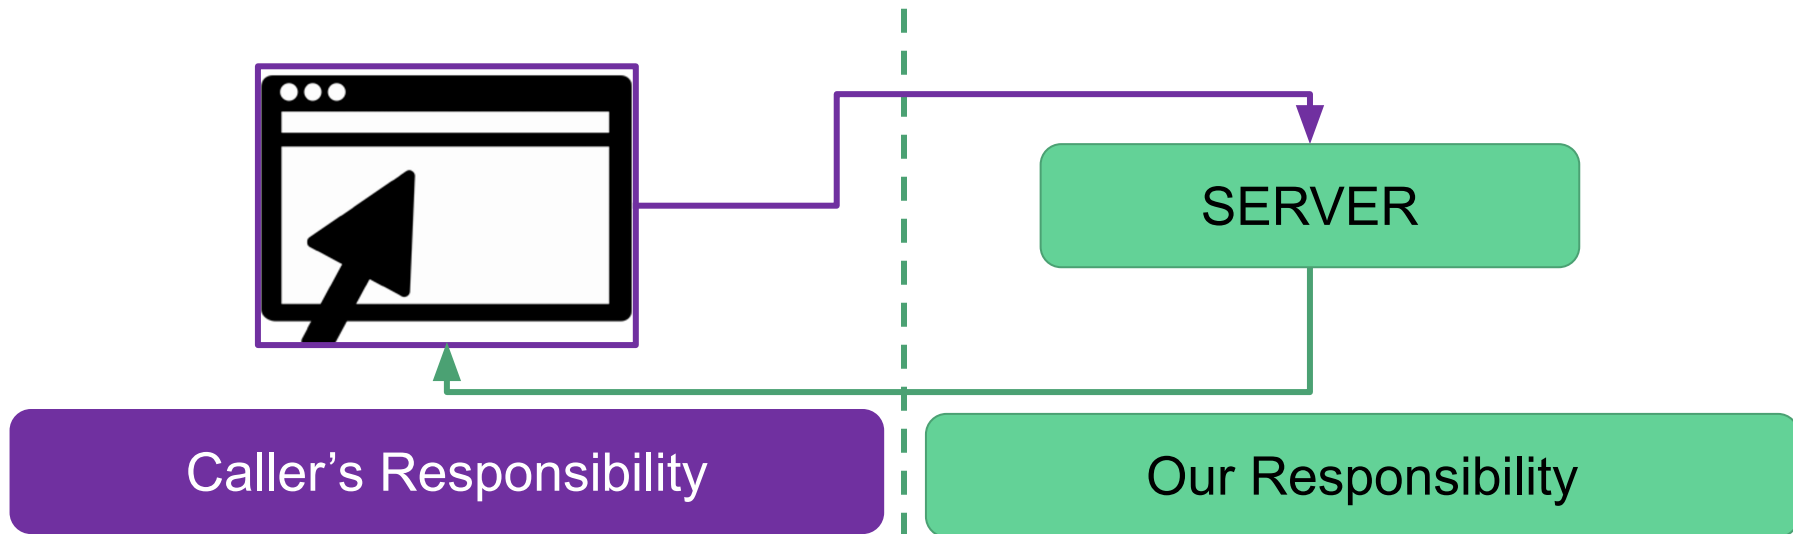
# Welcome to the jungle

**If your work needs to have an impact, it needs to RUN OUTSIDE YOUR LAPTOP:**

1. Your code can be **inspected, modified, understood** by others, typically your technical colleagues: you need to write clean, modular, testable code and make your pipeline fully reproducible.
2. Your model can be **trusted** by others, typically, other stakeholders, who may or may not be technical folks: you need to “make sure” the model behaved as designed before pushing it in front of end-users.
3. Predictions can be **consumed** by others, typically anybody with an internet connection: you need to expose your model as an endpoint which returns predictions when supplied with the appropriate parameters.


## Part 3: Serving predictions

- If our model stays on our laptop, nobody will be able to use it!
- **Client-server architecture:** our model interacts with *many* remote clients through an API (also called “endpoint”) - we abstract away model code (and complexity) and expose a pure input-output interface: clients send us the input, we return a prediction.



## Part 3: Serving predictions

- If our model stays on our laptop, nobody will be able to use it!
- **Client-server architecture:** our model interacts with *many* remote clients through an API (also called “endpoint”) - we abstract away model code (and complexity) and expose a pure input-output interface: clients send us the input, we return a prediction.



HTML (+ CSS) + Javascript

SERVER

Python (Flask, FastAPI etc.)

Show me first!

# Intro to Flask applications

- We will be using Flask as a simple framework to serve model predictions after training
- Flask has several attractive features:
  - Helps with structuring both the front-end (the web page) and back-end (the endpoint)
  - Pure Python back-end
  - Minimal syntax for routing, GET / POST etc.



# Step 1: prepare a web page

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4   <title>{{ project }} app</title>
5   <script src="https://ajax.googleapis.com/ajax/libs/jquery/3.5.1/jquery.min.js"></script>
6 </head>
7 <script type="text/javascript">
8
9   $(function() {
10     $('#predict').click(function() {
11       event.preventDefault();
12       var form_data = new FormData($('#myform')[0]);
13       console.log(form_data);
14       $.ajax({
15         type: 'POST',
16         url: '/',
17         data: form_data,
18         contentType: false,
19         processData: false,
20       }).done(function(data, textStatus, jqXHR){
21         $('#result').text(data);
22       }).fail(function(data){
23         alert('error!');
24       });
25     });
26   });
27
28 </script>
29 <body>
30   <h1>{{ project }}</h1>
```

- Prepare a simple HTML page for users to interact with our endpoint.
  - It is not much different than the streamlit app we built before!
- Note that we use a simple Javascript function with jQuery to perform a POST request.

## Step 2: prepare the Flask back-end application

```
32 # We need to initialise the Flask object to run the flask app
33 # By assigning parameters as static folder name, templates folder name
34 app = Flask(__name__, static_folder='static', template_folder='templates')
```

- Initialize a Flask app in [app.py](#)
  - Note the *templates* folder contains the HTML we created before!
- Make sure the app is started when we run “flask run”: the script will spin up a web server that will be ready to listen for incoming requests (from our HTML page, of course)

```
54 if __name__=='__main__':
55     # Run the Flask app to run the server
56     app.run(debug=True)
```

## Step 3: load the ML model in memory

```
13 ##### THIS IS GLOBAL, SO OBJECTS LIKE THE MODEL CAN BE RE-USED ACROSS REQUESTS #####
14
15 FLOW_NAME = 'MyRegressionFlow' # name of the target class that generated the model
16 # Set the metadata provider as the src folder in the project,
17 # which should contains /.metaflow
18 metadata('../src')
19 # Fetch currently configured metadata provider to check it's local!
20 print(get_metadata())
21
22 def get_latest_successful_run(flow_name: str):
23     "Gets the latest successfull run."
24     for r in Flow(flow_name).runs():
25         if r.successful:
26             return r
27
28 # get artifacts from latest run, using Metaflow Client API
29 latest_run = get_latest_successful_run(FLOW_NAME)
30 latest_model = latest_run.data.model
```

- As in all Python scripts, what is declared outside the functions is “global”.
  - In this context, this means that objects can live *across requests*: if the client ask for prediction 1 and then prediction 2, we do not need to reload the model, as it is already **in memory!**
- We want the model to be “global” as retrieving the model from Metaflow storage may be slow and expensive.


## Step 4: defining our endpoint

```
36 @app.route('/', methods=['POST', 'GET'])
37 def main():
38
39     # on GET we display the page
40     if request.method=='GET':
41         return render_template('index.html', project=FLOW_NAME)
42     # on POST we make a prediction over the input text supplied by the user
43     if request.method=='POST':
44         # debug
45         # print(request.form.keys())
46         _x = request.form['_x']
47         val = latest_model.predict([[float(_x)]])
48         # debug
49         print(_x, val)
50         # Returning the response to the client
51         return "Predicted Y is {}".format(val)
```

- We use a decorator to define our route (empty in this case, but could be, say, “predict”).
  - If it was “predict”, our server would be listening for calls at **URL/predict**
- We distinguish between page load (GET) and request from a prediction (POST).

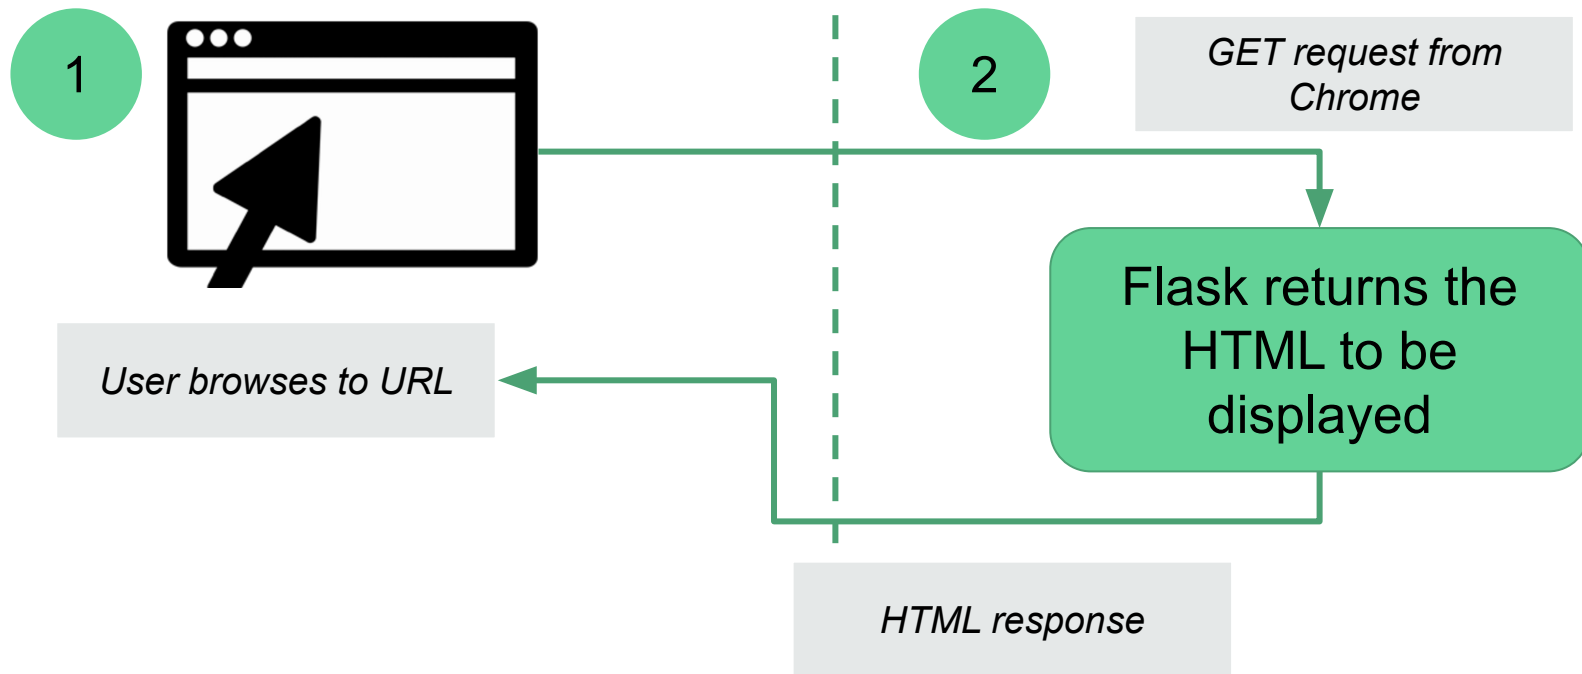
# The full workflow

1




*User browses to URL*

# The full workflow



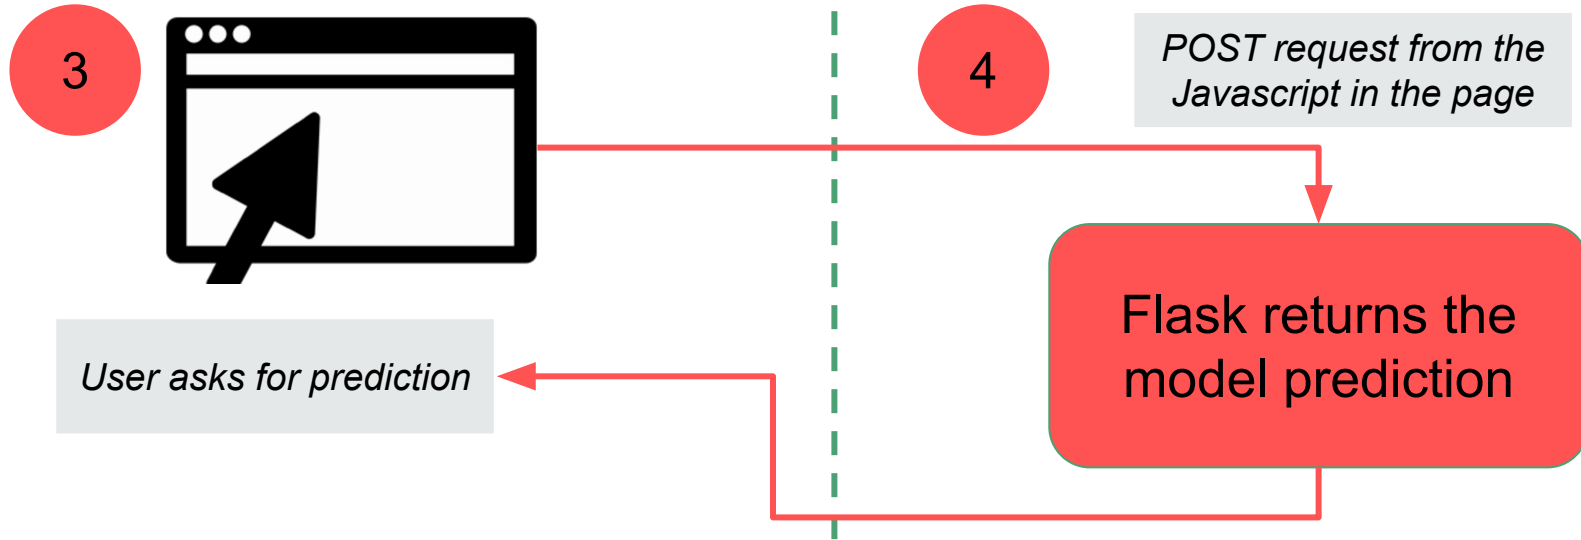
# The full workflow

3



*User asks for prediction*

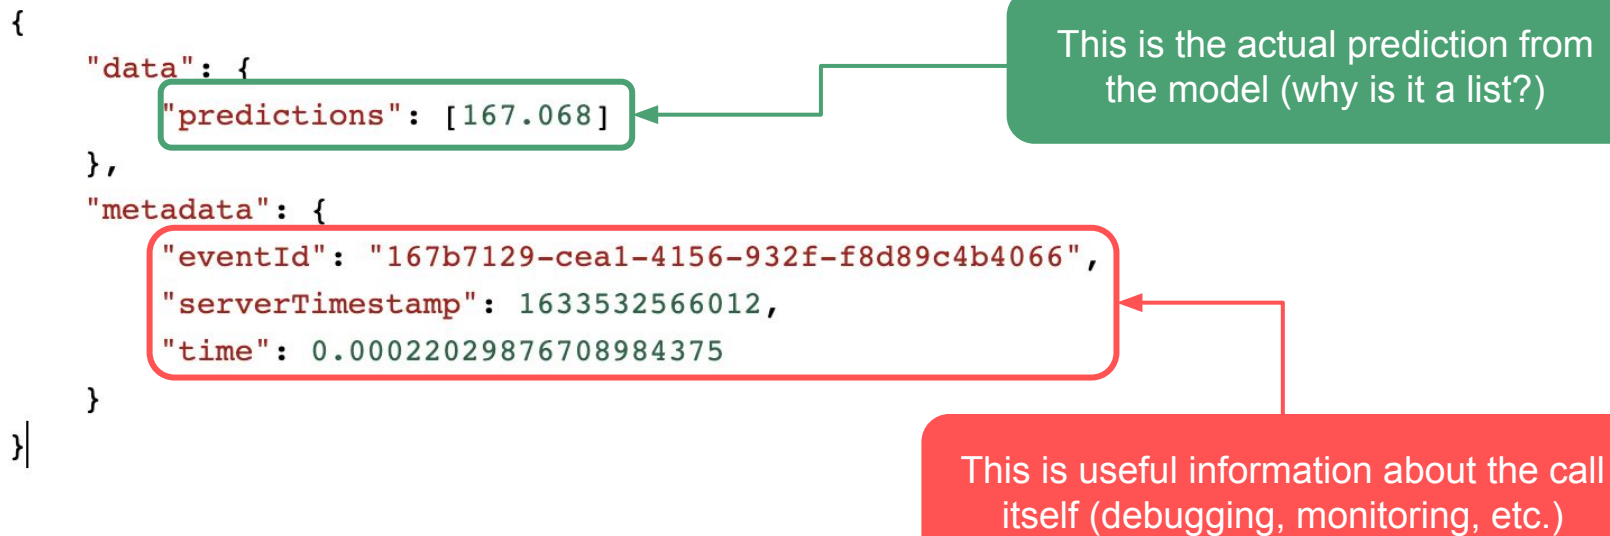
# The full workflow



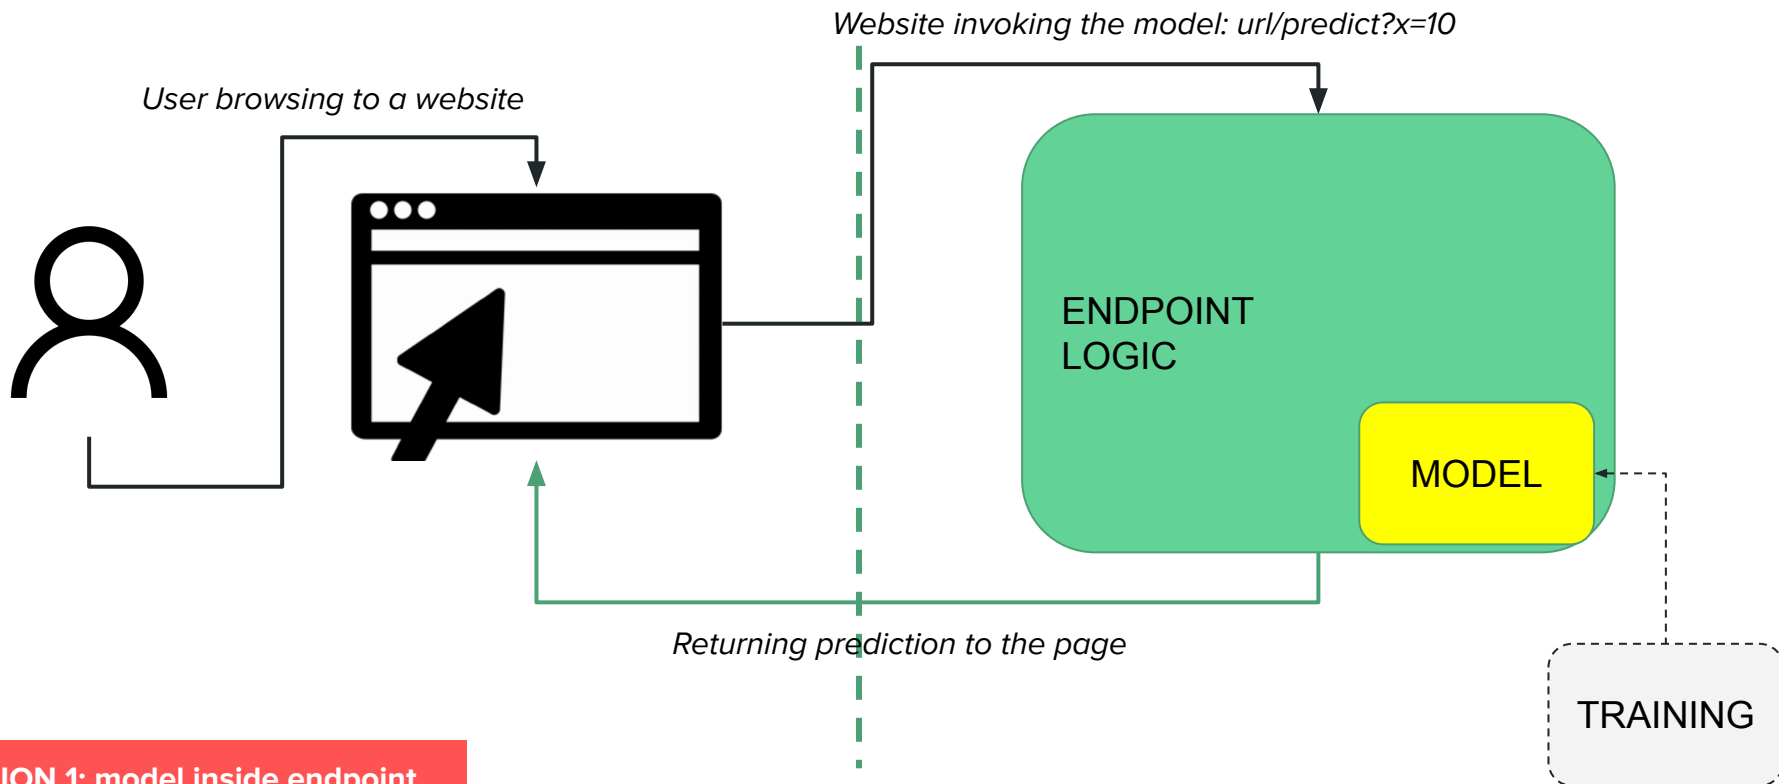


## BONUS: Structuring the response

- Now **everybody** with the URL can use your awesome model!
- Can we make the response a bit clearer?

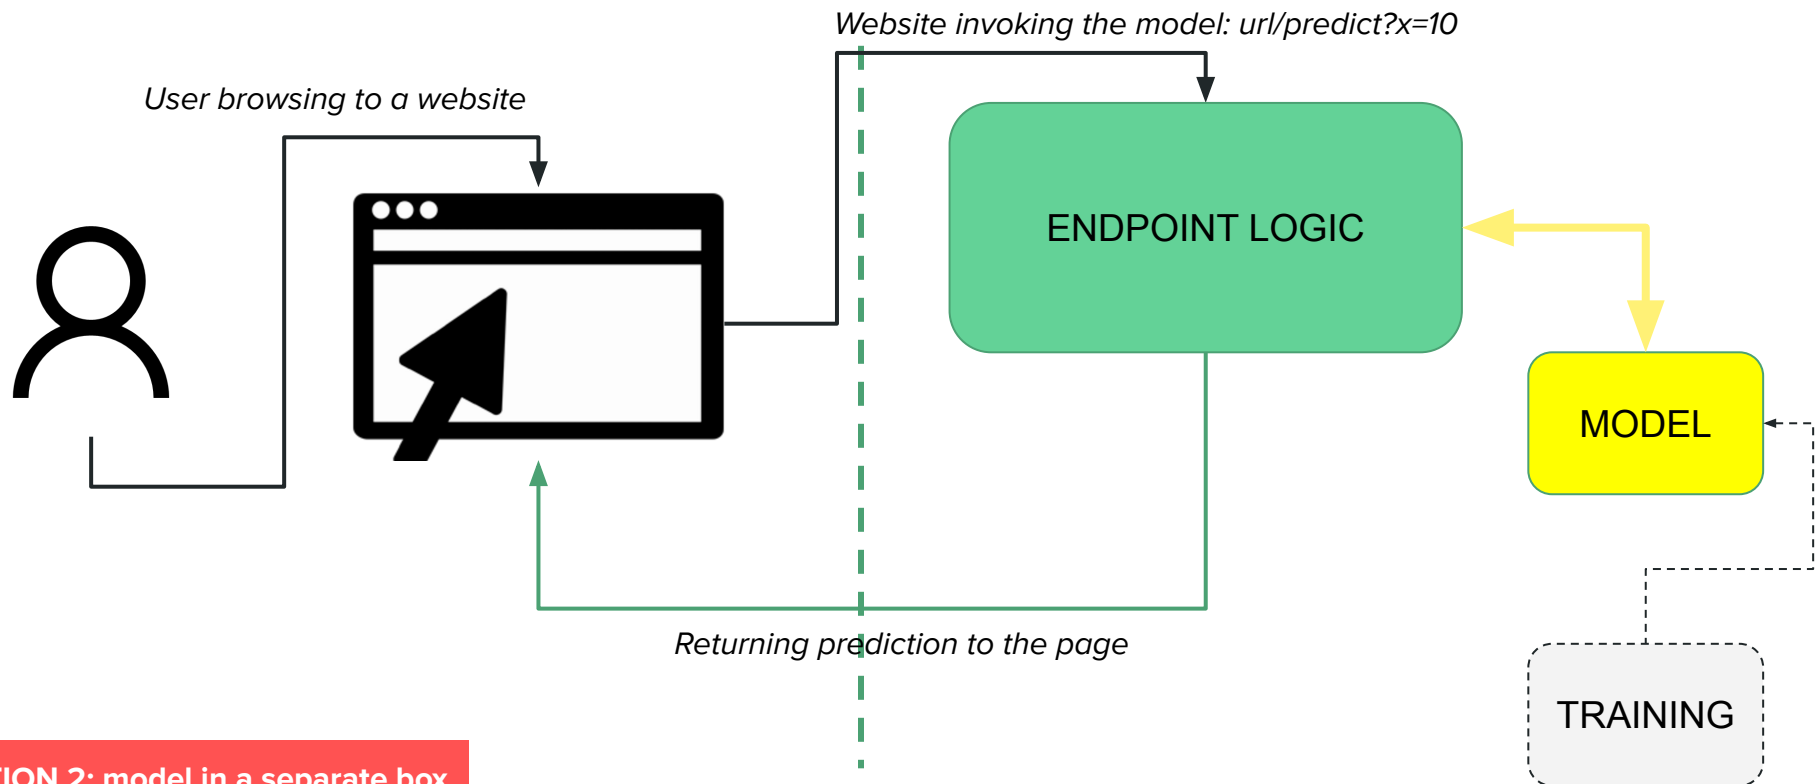


# Scenario 1: Endpoint with Model (ours)



**OPTION 1: model inside endpoint**

## Scenario 2: Endpoint + Model



**OPTION 2: model in a separate box**

# Flask... in the cloud

## Deploying ML models To the Web with Flask on AWS EC2 Instance



- There's 1M and one tutorials on how to use the very same tools (a Flask web app, a simple HTML + Javascript page) to port your app into the cloud.
- **BONUS points for your final demo if you can show your endpoint in the cloud**, either through an EC2 or Streamlit Cloud (**make sure to use the free resources / plan to avoid incurring in costs!**)

# Alternative deployment scenarios

There is a ton of alternatives when it comes to *serving predictions* from the cloud, ranging from pure infrastructure to fully managed services. For example:

- You can deploy your model manually on a virtual machine, by installing Flask and run through screen (like they do [here](#))
- You can deploy your model through a web app hosted by Elasticbeanstalk (like they do [here](#))
- You can deploy your model through a web app hosted by Fargate (like they do [here](#))
- You can deploy your model through Sagemaker, and expose it through a lambda (like [we did in the 2021 repository](#))

# After deployment: monitoring

We are not going to discuss monitoring, as we are not launching new apps in this course (for now!). However, after our model is live we need to:

- monitor how the pipeline is doing:
  - How is the new data coming in?
  - Does the model need re-training?
  - Is my new model better than the old one?
- check what users are doing with it!

# After deployment: monitoring



We are not going to discuss monitoring, as we are not launching new apps in this course (for now!). However, after our model is live we need to:

- monitor how the pipeline is doing:
  - How is the new data coming in?
  - Does the model need re-training?
  - Is my new model better than the old one?
- check what users are doing with it!
  - You never know how people would use stuff!



# The adventure never stops!

There is a ton of recent developments in the “MLOps” space (we do our small part as well in the community). If you want to know more, reach out!



## no-ops-machine-learning

A PaaS End-to-End ML Setup with Metaflow, Serverless and SageMaker.

This repo is an end-to-end model building exercise, showing how to go from a dataset file to an endpoint serving predictions through a repeatable, modular and scalable process (DAG-based). More importantly, all of this is achieved in minutes from a developer laptop, without explicitly deploying/maintaining *any infrastructure*.

For the full context and details about the tutorial, please refer to the [blog post](#).

### Overview

### README.md

## You Don't Need a Bigger Boat

An end-to-end (Metaflow-based) implementation of an intent...

This is a WIP - check back often for updates.

### Philosophical Motivations

There is plenty of tutorials and blog posts around the Internet on data pipelines and tooling. However:

- they (for good pedagogical reasons) tend to focus on *one tool / step* at a time, leaving us to wonder how the rest of the pipeline works;
- they (for good pedagogical reasons) tend to work in a *toy-world* fashion, leaving us to wonder what would happen when a *real dataset* and a *real-world problem* enter the scene.